

Cytochrome P450 Inhibitor Classification with Statistical Learning

D.Siegle

Abstract

This project compares methods of Cytochrome P450 inhibitor prediction based on compound structures. CYPs are a natural first choice in which to develop *in silico* models because of their central role in drug candidate rejection due to adverse drug-drug interactions. Several non-linear, high-dimensional classification models were built and compared using a large, publicly available high throughput screening luminescence assay (PubChem AID1851) against five CYP isozymes (1A2, 2C9, 2C19, 2D6 and 3A4). The methods compared are a Bayesian binary QSAR method from the Molecular Operating Environment, and 3 standard machine learning methods implemented in the Python programming language; κ -Nearest Neighbor, Random Forests, and Support Vector Machines. They all performed well, with the methods implemented in freely available software performing as well or better than the one in software that is widely accepted in industry.

Contents

1	Background	7
1.1	Challenges to Drug Discovery Productivity	7
1.2	Cytochrome P450 Superfamily	10
1.3	Rationale for early compound profiling <i>in silico</i>	11
1.4	Quantitative Structure-Activity Relationship	14
1.5	Statistical Machine Learning	18
1.6	should I mention high throughput or -omics tech as a source of data here	19
1.7	Reproducibility	19
1.8	Aims	21
2	Materials and Methods	23
2.1	Review of PubChem Assay 1851	23
2.2	Data set preparation	26
2.2.1	Molecular descriptor generation	27
2.3	Feature selection	28
2.4	Modeling	29
2.4.1	Methods of Correlation of Descriptors to Activity Data	30
2.5	Model Evaluation and Validation	33
3	Results	40
3.1	2C19 MOE models	42
3.2	2C9 MOE Models	44
3.3	1A2 MOE Models	45
3.4	2D6 MOE Models	46
3.5	3A4 MOE Models	48
3.6	Comparison of Methods	51
3.6.1	1A2	51

3.6.2	2C9	51
3.6.3	2C19	51
3.6.4	2D6	51
3.6.5	3A4	51
4	Discussion	55
4.1	Other attempts at modelling Assay 1851	55
4.1.1	Cheng's results compared to mine	56
4.2	Sources of error	56
4.3	Machine Learning	57
5	Conclusion	60

List of Figures

1.1	Typical QSAR Workflow	22
2.1	Supervised Learning Scheme for Clasification	24
2.2	Data download and preparation	27
2.3	Isozyme data set preparation	28
2.4	Supervised Classification Model Training and Validation Work- flow	30
2.5	General Data Analysis Workflow	37
3.1	2C19 MOE Model Accuracy	43
3.2	2C19 MOE Model Accuracy	45
3.3	1A2 MOE Model Accuracy	47
3.4	2D6 MOE Model Accuracy	49
3.5	3A4 MOE Model Accuracy	50
3.6	Inhibitor Classification Method Comparison on the Training Set	53
3.7	Inhibitor Classification Method Comparison on the Test Set .	54

List of Tables

3.1	Results from MOE binary classification	40
3.2	2C19 MOE Model Accuracy	42
3.3	2C9 MOE model results	44
3.4	1A2 MOE model results	46
3.5	2D6 MOE model results	48
3.6	3A4 MOE model results	48
3.7	Comparison of Classification Methods for 1A2	51
3.8	Comparison of Classification Methods for 2C9	51
3.9	Comparison of Classification Methods for 2C19	52
3.10	Comparison of Classification Methods for 2D6	52
3.11	Comparison of Classification Methods for 3A4	52

Acknowledgements

I want to thank ...

Chapter 1

Background

1.1 Challenges to Drug Discovery Productivity

The goal of pharmaceutical sciences is to identify and market safe and efficacious drugs. Toxicity is a large contributor to drug candidate attrition in drug development. Inhibition of enzymes in the cytochrome P450 superfamily is a major source of toxicity in human and animal models because of their role in first pass metabolism. If a compound inhibits cytochrome P450 it is likely to lead to toxic effects when first pass metabolism fails to clear other therapeutic compounds or alters their pharmacokinetics. Because the discovery and development of a each new pharamceutical drug requires more than a billion dollars across an average of 12 years, it would be useful to know whether a compound of interest inhibits cytochrome P450s as early as

possible.

The Cytochrome P450 superfamily is large and varied. Different isozymes metabolize different substrates. Even within and between individuals pharmacokinetic and pharmacodynamic variability can be high for reasons not entirely characterized. Designing assays for every isozyme and SNP variant and running them against every compound of interest as a general strategy is cost prohibitive. This project was designed to test the predictive power of computational methods for cytochrome P450 inhibition potential based simply on knowledge of chemical structure.

There has been an upward trend in drug development costs for three decades for a variety of reasons. Integrated computational approaches have been proposed as one way to control costs in pharmaceutical research and development. The foundational skills demonstrated in this thesis are needed to pursue a systems approach to drug development that industry has recently turned toward as a way to boost R&D productivity. [Visser et al., 2014, Berg, 2014] It is also part of the larger project of *in silico* drug-development, attempting to reduce reliance on exploratory *in vivo* and clinical drug testing and increase the number of effective treatments for patients and decrease the amount of time to develop them.

Math is the new microscope [Cohen, 2004]

There approach of associating compound structure with bioactivity goes back at least to Hansch [Hansch and Fujita, 1964] and has progressed under the moniker of Quantitative Structure Activity Relationship (QSAR). QSAR

started with direct measures of chemical compounds and later derived features and used them to then build expert systems or statistical models that tried to predict biological activity. During those same decades the field of machine learning emerged; using computation in an analogous and more general way – to associate features with results, inputs with outputs. Machine learning can be thought of as using algorithms to figure out how to perform important tasks by generalizing from examples. These algorithms are of course usually laborious, which necessitates their execution by computer.

Statistical machine learning builds upon the peer prediction of machine learning. It also allows for prediction but focuses more on models and methods that can be used by scientists and engineers. [James et al., 2013] Further extension of statistical learning to the pharmaceutical sciences can lead to important contributions to systems pharmacology. Or rather statistical learning methods are an important precursor to the needs of systems biology and systems pharmacology modeling.

This project compares different techniques for Cytochrome P450 inhibition prediction in the framework of statistical learning, for the five isozymes of Cytochrome P450 that are involved in metabolism of 90% of all therapeutic drugs.

Because replicability is one of the main principles of the scientific method. As much as possible, the code and data used in this study is publicly available and source controlled. Reproducibility as a practice is a habit, and a good one to get into. Every attempt was made to make the materials and methods

for this project reproducible in a completely automated way to allow for validation of results and extension of the work.

1.2 Cytochrome P450 Superfamily

One of the largest and most functionally diverse protein superfamilies is the cytochrome P450 family of hemoproteins. From bacteria to humans, the functional breadth of cytochrome P450 activity is far ranging. At the latest count there were significantly more than 2000 identified cytochrome P450 genomic and cDNA sequences that have been divided into a total of 265 different families. [Danielson, 2002] Cytochromes P450 appear in every kingdom from bacteria to higher eukaryotes. Multiple cytochrome P450 genes can be expressed simultaneously as different isozymes and the number of genes per species is highly variable with a tendency for higher eukaryotes to possess more. The cytochromes P450 (CYPs) constitute the major enzyme family capable of catalyzing the oxidative biotransformation of most drugs and other lipophilic xenobiotics and are therefore of particular relevance for clinical pharmacology. This central role that these ubiquitous proteins play as phase I enzymes in human drug metabolism makes them very important to the pharmaceutical industry.

CYP families are classified based on pairwise amino acid sequence identity among individual members. Families CYP 1-3 are involved in phase I metabolism of human drugs and xenobiotic compounds, whereas other CYP

families (CYP 4, 11, 17, 17 and 21) are involved in the metabolism of endogenous compounds such as fatty acids, steroids, eicosanoids, bile acids and fat soluble vitamins[Singh et al., 2011]

There is a considerable substrate overlap between enzymes of this superfamily. Being broadly specific with respect to their substrates, CYPs are therefore susceptible to inhibition by a large variety of chemical compounds. It follows that the CYP enzymes that are involved in the oxidative metabolism of drugs play a major part in the activation and elimination of therapeutic drug molecules. CYP inhibition leading to decreased elimination and/or changed metabolic pathways of other substrates is a major cause of adverse drug-drug interactions. [Lapins et al., 2013] So adverse side effects of drug-drug interactions are an important area of inquiry, especially during the research phase of drug discovery.

1.3 Rationale for early compound profiling *in silico*

Drug discovery is a multi-parameter optimization process in which compounds are optimized for interaction with their target while minimizing off-target activities. [Zlokarnik et al., 2005] Adding drug-like properties to those compounds may minimize the risk of making potent and target-class selective compounds in a chemical space that is biologically inaccessible but at the same time does little to address the complexities of specific drug-drug

interactions.

Because of the complex nature of toxicity, safety prediction is considered more challenging than efficacy prediction. Toxicity mechanisms may be unknown or poorly characterized in higher organisms, and similar pathways and targets may be associated with different toxicities and adverse events. Toxicity prediction must also encompass a number of complex interactions and remain alert to the possibility of finding the unexpected. For instance, toxicities could result from on-target effects due to incomplete knowledge or inadequate target validation, or from off-target effects mediated via unknown molecules and mechanisms, or even from genetic variation or comorbidities in any of the previously mentioned pathways. [Kruhlak et al., 2012]

Techniques for high-throughput *in vitro* screening of CYP inhibition have been developed and implemented broadly in drug discovery pipelines across pharmaceutical companies and research institutions resulting in the generation of large datasets. Some of this accumulation of data has been released through academic research initiatives (e.g. PubChem Bioassays AID 410 and 1851). These collections enable development of structure-activity relationship models for *in silico* prediction of CYP inhibition by a much larger pool of researchers than those who designed and carried out the assays.

The promise of *In silico* screening that remains very appealing is that, with a steady increase in computing power, screening cost could become negligible. The hope is that virtual compounds could be screened for CYP liabilities in order to realize savings and reduce the number of candidates with ques-

tionable prospects that a team would otherwise synthesize.[Zlokarnik et al., 2005] Achieving this would also turn the costly animal toxicity test phase of pre-clinical research into a validation step rather than a screening step.

The ability to predict clinical safety based on chemical structures is also becoming an increasingly important part of regulatory decision-making. QSAR models are currently used by industry and by regulators to evaluate safety concerns and possible nonclinical effects of a drug when adequate safety data is absent or inconclusive.[Kruhlak et al., 2012]

Many high-throughput technologies are now available to detect P450 inhibitors. High-throughput screening data can be used to guide medicinal chemists away from these interactions in an early stage. In certain cases it might also identify the inhibition issue by targeted modification of the CYP interacting functionality. To be generally useful, P450 inhibition screens need to be calibrated against standard methods and preferably also tested with a large set of drugs, for which human drug-drug interaction outcome is known. [Zlokarnik et al., 2005] This should decrease the number of withdrawals of novel drugs from the market due to inhibition of major P450 isozymes.

The United States is about to become the last country to accept QSAR in the drug approval process. The drafting of International Committee on Harmonization (ICH) M7 guideline can be viewed as setting a precedent for possible future, broader regulatory applications of QSAR modeling. ICH M7, will for the first time specify that – under very specific conditions – the results of QSAR computational toxicology predictions will be considered

sufficient for genotoxic contaminants of pharmaceuticals under consideration and eliminate the need for laboratory testing. [Kruhlak et al., 2012]

1.4 Quantitative Structure-Activity Relationship

Quantitative structure-activity relationship (QSAR) modeling is generally accepted as the construction of predictive models of biological activities as a function of structural and molecular information of a compound or compound library.[Nantasenamat, 2009]

QSAR models describe the correlation between molecular features and activity at a given end point of interest. There have been attempts to make structure activity models (SAR) constructed by using human expert knowledge (expert rule-based), QSAR models are typically defined as those that use statistical methods to analyze the mathematical correlations between molecular features and activity.

It follows from the 'similarity principle' that new and untested compounds possessing similar molecular features as known compounds are assumed to possess similar activities and properties. In this way, QSAR models can make it possible to predict the biological activities of a given compound as a function of its molecular structure. Several successful models have been published over the years which encompass a wide span of biological and physicochemical properties.

A QSAR model defines the mathematical relationships between descriptors and biological activities of known molecules. This differs from receptor binding-based efficacy prediction by taking into account binding site characteristics as well as molecular docking analysis. In contrast, receptor-binding methods attempt to predict drug efficacy based on known mechanisms of action and medicinal chemistry by individually studying molecular interactions between a drug and targets/receptors. [Kruhlak et al., 2012]

Applied QSAR, as described above, has typically been used for drug discovery and development but has also been used to correlate molecular information with other physiochemical properties. This latter approach is termed quantitative structure-property relationship (QSPR). Derived molecular parameters can account for hydrophobicity, topology, electronic properties, and steric effects. These characteristics of compounds can either be determined empirically through experimentation or theoretically via computational chemistry as needed. [Nantasenamat, 2009] The parameters derived from compound structure are referred to as molecular descriptors.

Molecular Descriptors

Molecular descriptors can be thought of as the mathematical representation of essential information of a molecule in terms of its own physiochemical properties. Depending on the needs of the analysis, properties considered can be electronic, geometric, hydrophobic, constitutional, lipophilic, steric, solubility, quantum chemical, or topological. From a practical viewpoint,

molecular descriptors are chemical information that is encoded within the molecular structures. [Nantasenamat, 2009]

Molecular features can be either experimentally measured or calculated values. They come in the form of simple physiochemical properties such as logP or logarithmic acid dissociation constant (pKa), numerical representations of substructure fragments, or purely mathematical descriptors. Mathematical descriptors are chemical structural features represented in numerical form, and range from simple atom counts to the product of complex equations that describe electron distribution across a molecule. [Kruhlak et al., 2012]

Molecular descriptors as predictors in QSAR modeling are typically less precise than the lock and key relationships that underpin the docking approach to computer-aided drug design. The basic assumption in QSAR modeling is that similar molecules exhibit similar biological activity so that physiochemical properties and/or structural properties of a molecule encoded as molecular descriptors can be used to predict the biological activity of structurally related compounds with some degree of confidence.

Modeling

QSAR models can be described as global or local. Global models incorporate chemicals with a range of molecular features acting across the spectrum of chemical pathways, whereas local models are highly focused on a single chemical class and end point. Although local models generally have much higher accuracy, their narrow domain of applicability renders them imprac-

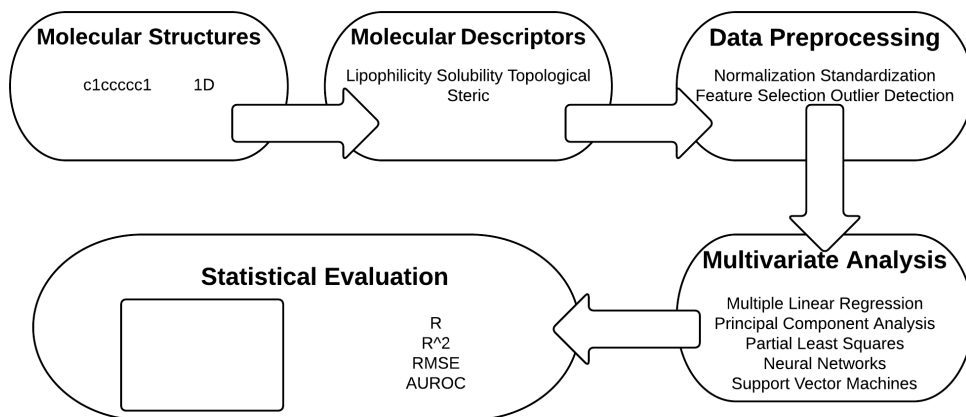
tical in most regulatory environments where predictions need to be made for structural classes covering active pharmaceutical ingredients to metabolites, degradants, reagents, and synthetic intermediates. [Kruhlak et al., 2012]

The construction of QSAR models typically follows two main steps:

- Description of molecular structure and derivation of descriptors
- Multivariate analysis correlating molecular descriptors with observed activities.

Additional intermediate steps are also crucial for successful development of such QSAR models and include data preprocessing and statistical evaluation. [Nantasenamat, 2009]

Figure 1.1: Typical QSAR Workflow



A typical QSAR workflow involves chemical structure management, descriptor calculation, and statistical analyses that are treated as separate steps

and often performed by non-integrated software packages. This can lead to low throughput and even the lack of possibility of performing predictions for new compounds or updating the models when new data become available, depending on the workflow.

According to Kruhlak, et al., the successful development of a QSAR model for safety prediction requires a sufficient amount of high-quality data, the appropriate selection of descriptors, the availability of one or more suitable statistical or mathematical models and an effective training and validation strategy. [Kruhlak et al., 2012]

1.5 Statistical Machine Learning

Machine learning algorithms figure out how to perform important tasks by generalizing from examples. A concise definition of statistics is as the applied science that constructs and studies techniques for data analysis (Jan de Leeuw) Statistical learning refers to a set of approaches for estimating a function that describes a dataset as a precursor for prediction or inference.[James et al., 2013] Statistical machine learning constructs that function by generalizing from examples, i.e. data.

Leo Breiman wrote a landmark paper that documented the beginnings of this approach. He said 'There are two cultures in the use of statistical modeling to reach conclusions from data. One assumes that the data are generated by a given stochastic data model. The other uses algorithmic

models and treats the data mechanisms as unknown’ [Breiman, 2001] He also claims that the statistical community had traditionally preferred the first view.

Classification is a well understood area of machine learning. A classifier is a system of inputs, typically a vector of discrete and/or continuous feature values which outputs a single discrete value, the class.[Domingos, 2012]

The current state of machine learning is fundamentally a subset of optimization and has found its biggest successes in fields where there are far more variables than parameters. The ultimate goal of machine learning in this case is to generalize beyond examples in the training set, because no matter how much data we have, it is unlikely that we will see those exact examples again at test time.[Domingos, 2012]

In the context of QSAR then, with enough prior knowledge, machine learning may be a practical approach to fill in the gaps of clinical knowledge for any relevant CYP isozyme queried against any known compound of interest so long as the compound structure is specified.

This is a wider view than traditionally encountered in most applied science education where simple hypothesis testing predominates. Opening up the field of data analysis like this brings opportunity for exploration but also new concerns, such as the ’curse of dimensionality’, ’degrees of freedom of the analyst’, ’black box algorithms’, difficulty of bias estimation and the ’no free lunch theorem’.[Boulesteix and Schmid, 2014]

1.6 should I mention high throughput or -omics tech as a source of data here

1.7 Reproducibility

Science conducted in an open fashion confers the following benefits

- Reproducibility of experiments allows other researchers to use the exact methods to calculate the relations between biological data.
- Faster development of disease models and therapeutic treatments due to the reuse of existing knowledge. Projects can be built upon existing results more easily or extend the research in directions unanticipated by the original team. First-pass results can be subject to new analysis, and a second look at compounds with interesting side effects can lead to serendipitous discoveries
- Increased quality as a result of having more researchers studying the same topic to provide a layer of assurance that errors will not propagate.
- Long-term availability of data and code. If these resources are not tied to businesses or patents, then they can be posted to multiple repositories to ensure that they are available in the future. [Prli and Lapp, 2012]

High profile case of inability to replicate the effects of major cancer drug findings. [] We don't know which drugs because those findings are not public

either.

Information generation can happen far faster and is much more common than data analysis and knowledge creation in the biological sciences.

There are also organizational issues to overcome. More biologists trained in quantitative and statistical methods for analyzing large data sets are needed.

Large -omics data collections can be expensive to generate and also to manage. Efforts are often distributed and decentralized, leading to duplications of effort and lack of common standards. For pharmaceutical drug discovery researchers, the long-term commitment and infrastructure required to support these approaches is not compatible with current restructuring trends in the industry.[Berg, 2014]

New research findings, supporting data and methods should, therefore, be made publicly available for independent verification and replication in order not to delay medical advances.

1.8 Aims

The aim of this project is to build QSAR models that quantify the risk of off target effects for candidate molecules by building statistical machine learning models that make binary classification of compounds based on their structure, as inhibitors or non-inhibitors of Cytochrome P450.

We will also compare a well accepted, commercial method of binary clas-

sification with a three open source implementations of QSAR model building according to the following plan:

- Build Binary QSAR models in the Molecular Operating Environment (MOE).
- Develop and implement comparable methods in open source software.
- Evaluate and compare results from all models.
- Perform this analysis as reproducible research.

Chapter 2

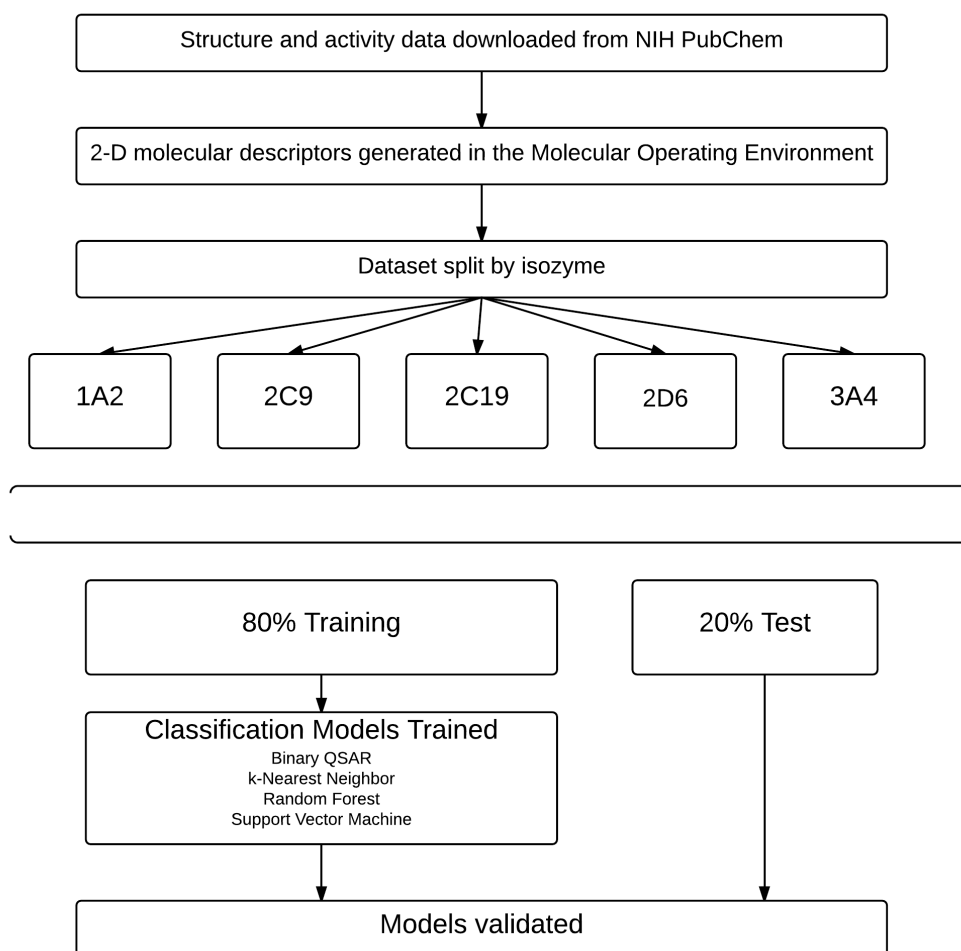
Materials and Methods

2.1 Review of PubChem Assay 1851

PubChem BioAssay 1851 contains data for inhibition of five major CYP isoforms (CYP1A2, CYP2C9, CYP2C19, CYP2D6 and CYP3A4) by 17,143 chemical compounds. [Veith et al., 2009] The the tested compounds were all drugs or drug-like compounds. The chemical space revealed that the majority of compounds had molecular weight below 500 daltons and logP below 5. [Lapins et al., 2013]

The assay used low low-fluorescence substrates, which are converted to more-fluorescent metabolites. The progression of the reaction is measured by an increase in fluorescence intensity during CYP metabolism of the substrate. Inhibitors of a particular CYP reduce the rate of metabolism of the substrate and this results in a decreased fluorescence signal. [Zlokarnik et al., 2005]

Figure 2.1: Supervised Learning Scheme for Classification



The most recent technology developed for CYP inhibition is based on substrates that release luciferin as the metabolite. This is a coupled assay system in which addition of luciferase and ATP converts the freed luciferin to des-carboxyluciferin with light emission. The format is similar to fluorescence methods, requiring addition-only manipulations and employing luminescence plate readers. [Zlokarnik et al., 2005]

The assay obtained luminescence readings at a range of compound concentrations. Here I need a paragraph describing the Hill equation and its suitability for curve fitting

In the dataset, compounds are classified as active or inactive inhibitors for each CYP with an activity cutoff set to $AC_{50} = 10\mu M$ (AC_{50} , activity concentration 50, refers to the concentration that is required to elicit half-maximal effect). However, in cases where the dose-response curve for a compound showed poor fit or the inhibition efficacy was below 60%, the assay results were regarded as inconclusive. [Lapins et al., 2013]

Compounds were characterized by their Activity Score and regarded as inhibitors if their activity score ranged between 40 and 100. PubChem Activity Score is assigned based on an AC_{50} values, which was combined with a confidence measure. Combining measures for completeness of a dose-response curve and efficacy of inhibition, resulted in the Activity Score, where a larger value indicates higher inhibitory activity and/or higher confidence in inhibitory assay result. Compounds with an activity score equal to zero are considered as non-inhibitors while compounds with activity scores above 0

and up to 40 are considered inconclusive. [Lapins et al., 2013]

The data from Pubchem Assay 1851 is available for download from the NIH website. The interface changes from time to time, but I downloaded two files which comprised the entire dataset for the experiment. The first file, a structures file, contained the structural information encoded in the simplified molecular-input line-entry system (SMILES) format for each tested compound with corresponding Structure ID (SID) and Compound ID (CID) as assigned by NCBI. Another file, also organized by SID and CID, contained all of the luminescent responses from the high-throughput screen and the fitted parameters that are summarized by the activity score.

2.2 Data set preparation

Both data files from Bioassay 1851 were downloaded as comma separated value (.csv) files and merged together based on the SID column using functions from Python’s pandas library, which is made for handling, manipulating and reshaping structured data.

The preliminary steps of data preprocessing typically requires data cleaning as raw data often contain anomalies, errors, or inconsistencies such as missing data, incomplete data, and invalid character values which may cause trouble in data analysis if left untreated. It is more complicated when data are collated from many formats that require harmonization and redundancy elimination. [Nantasenamat, 2009] In this case, minimal preprocessing was

required due to the polished nature of the source files.

Figure 2.2: Data download and preparation

2.2.1 Molecular descriptor generation

First the merged file was loaded into a database in the Molecular Operating Environment. MOE functionality was used to obtain the washed configuration of compounds by removing the salts and finding an energy-minimized conformation. MOE was then used to calculate descriptors based on the molecular structures. The entire suite of MOE 2-D descriptors was selected for descriptor generation, resulting in 186 additional columns of nominal, ordinal and continuous vales appended to the database. The resulting master file was saved in .csv format.

The dataset was then split by isozyme into 5 separate files that each contained only the SID, the Activity Score, and all 186 of the MOE 2-D descriptors using a script in the Python programming language.

For each isozyme, the number of active inhibitors was far outnumbered by the number of inactive inhibitors. The data file for each isozyme was subjected to a Python script that separated the inactives from actives. Then the order of the inactives was randomly shuffled and the column trimmed to the length of the activity column, thereby balancing the number of inactives and actives as required by some of the statistical methods.

Next the script randomly shuffled the balanced datasets and split them

Figure 2.3: Isozyme data set preparation

into a training set and a test set. The training set held 80% of the original values and the test set held 20% of the original values. The split was not based on activity score. The ratio of active/inactive in each split was inspected to check if they were still acceptably balanced.

For each use of randomness in computation, a seed number was set for the random number generator to ensure reproducible results.

The balanced and split datasets are saved to Figshare.com for permanent, free and open access. (<http://dx.doi.org/10.6084/m9.figshare.1181846> and <http://dx.doi.org/10.6084/m9.figshare.1066108>) All subsequent analyses use these same splits for comparability.

2.3 Feature selection

Typically data sets often contain redundant or noisy variables which make it more difficult for learning algorithms to discern meaningful patterns from the input. Such multicollinearity of the variables can either be used or treated if necessary to reduce computational resources required for model construction. [Nantasenamat, 2009]

There existed a great deal of variability in the range and distribution of each variable in the data set. However, this may pose a problem for algorithms using distance measurements in the learning step. These situations

are handled by applying statistical techniques such as min-max normalization or z-score standardiation.

In min-max normalization, the minimum and maximum value of each variable is adjusted to a uniform range between 0 and 1 according to the following equation: $X_{norm} = \frac{X - X_{min}}{X_{max} - X_{min}}$ In z-score standardization, essentially the variable of interest is subjected to statistical operation to achieve mean center and unit variance according to the following formula: $z = \frac{x - \mu}{\sigma}$ [Nantasenamat, 2009]

In situations where the data does not have a Gaussian (normal) distribution, simple mathematical functions can be applied to achieve normality or symmetry in the data distribution. A commonly used approach is to apply logarithmic transformation on the the variable of interest in order to achieve distribution approaching normality. This is typically performed on dependent variables such as the modeled biological/chemical properties of interest whereby IC50 may be transformed to logIC50 or -logIC50. Practically, such mathematical operation is applied to each individual value of a given variable of interest. [Nantasenamat, 2009]

2.4 Modeling

Balance number of positives and negatives

Split dataset into training and test set before looking at it.

Cross-validation

Molecular Operating Environment

SciPy ecosystem

2.4.1 Methods of Correlation of Descriptors to Activity Data

Figure 2.4: Supervised Classification Model Training and Validation Workflow

Binary QSAR in the Molecular Operating Environment

PLS regression is a quantitative method available in the Molecular Operating Environment that has shown poor predictive accuracy with this dataset as demonstrated by previous attempts in Dr. Zheng's lab.

The Binary QSAR approach outlined by Labute [Labute, 1999] and included in the MOE version, takes a Bayesian and probabilistic approach to classification of activity after reducing all the descriptors to their principal components.

To carry out this analysis, first load the training data into MOE and use a menu driven interface to initiate the Binary QSAR methods. A threshold value of 39 was selected; all activity values 40 and above will be considered active and all values 39 and below are considered inactive. The smoothing

parameter was left at the default value of 0.25. MOE automatically performs principal component analysis on high-dimensional datasets.

For each isozyme a number of models were built, each using a different number of principal components. Each principal component is orthogonal and uncorrelated to the rest, and each one captures a portion of the total variance inherent in the dataset. The assumption is that inclusion of more principal components leads to more of the variance being accounted for in a classification decision. However, since each principal component is a linear combination of all the variables, the benefits of dimensionality reduction comes at the cost of interpretability. For exploratory purposes, models with 2, 5, 10, 15, 20, 30, and 44 principal components were constructed.

MOE models were written to .fit files and the model report saved as a .txt file.

Modeling in Python

The Python programming language is a dynamically-typed, object-oriented interpreted language. Its primary strength lies in the ease with which it allows a programmer to rapidly prototype a project, although it also has a powerful and mature set of standard libraries that can facilitate large-scale production-level software engineering projects as well. Python has a very shallow learning curve and excellent online learning resources.

κ -Nearest Neighbor (κ NN)

The κ NN algorithm predicts the class of a test set object based on the class membership of its κ most similar training set objects. [Lapins et al., 2013]

Random Forest (RF)

Random Forest is a classifier that consists of multiple decision trees. A decision tree is made of nodes and branches. At each node the dataset is split based on the value of some attribute that is selected so that the instances of different classes are predominately moved to different branches.

Classification starts at the root node and is performed by passing the instances along the tree to leaf nodes. [Lapins et al., 2013]

To introduce diversity between the trees of a random forest, a small subset of all attributes is randomly selected to take decisions at each node of each tree. The classification decision is performed by considering results of all trees by majority vote. The optimal size of the forest and the number of attributes to consider at each node were found by performing five-fold cross-validation. [Lapins et al., 2013]

Support Vector Machines (SVM)

SVM is a machine learning technique for classification or regression that uses linear or non-linear kernel-functions to project the data into a high-dimensional feature space. [Lapins et al., 2013]

Correlation is performed in this hyperspace based on the structural risk

minimization principle i.e., aiming to increase the generalization ability of a model. [Lapins et al., 2013]

They applied the commonly used Gaussian radial basis function kernel optimal gamma (width of the kernel function) and error penalty parameter C were found after performing grid search on fivefold cross validation. [Lapins et al., 2013]

2.5 Model Evaluation and Validation

The building blocks of a successful QSAR model are the accuracy of the input data, selection of appropriate descriptors and statistical tools, and most importantly validation of the developed model. Validation is the process where reliability and relevance of a procedure are established for a pre-determined purpose. For QSAR models validation should target robustness, prediction performances and applicability domain (AD) of the models. [Lapins et al., 2013]

Statistical evaluation in QSAR modeling is essential to validate the model as well as to evaluate its predictive performance. The predictive performance of a data set can be assessed by dividing it into a training set and a testing set. The training set is used for constructing a model and then the predictive performance of that model is evaluated on the testing set. Internal performance is typically assessed from the predictive performance of the training set while external performance can be assessed from the predictive perfor-

mance of the independent test set that has never been seen by the training model. A commonly used approach for internal validation is known as N-fold cross-validation where a data set is partitioned into N number of folds. For example, in a 5-fold cross-validation 1 fold is left out as a testing set while the remaining 9 folds are used as the training set for model construction and then validated with that fold left out. In situations where the number of samples is limited, leave-one-out cross-validation is the preferred approach. In that case, the number of folds is equal to the number of samples present in the data set. [Nantasenamat, 2009]

In model building, 5-fold cross validation was used in conjunction with model building on the training set for all of the models built in Python. After the test set was applied, two statistical metrics were used to We used two statistical methods: the overall prediction accuracy and the Area Under the Receiver Operating Characteristic (AUROC) curve. We assessed the predictive ability of the models by performing cross-validation and external predictions.

Confusion Matrix

A confusion matrix (or error matrix or table of confusion) is a representational summary of the performance of a classifier. All data points used in model building are apportioned along one axis according to their actual class value, and along another axis according to the value predicted by that model. For binary classification this results in a square matrix with two 2 columns

and 2 rows.

		Prediction Class			
		p		n	total
Actual Class	p'	True Positive	TP	False Negative	FN P'
	n'	False Positive	FP	True Negative	TN N'
total		P		N	

In this study, results that are both actual inhibitors according to assay data and predicted inhibitors according to the model results are deemed True Positives (TP). Similarly, actual noninhibitors that are predicted to be so are called True Negatives (TN). In the case of actual inhibitors predicted to be noninhibitors, these are labelled False Negatives (FN) and are the result of Type I errors. Noninhibitors classified as inhibitors by the model are referred to as False Positives (FP) and are the result of Type II errors.

Accuracy

The accuracy metric provides general information about how many compounds are misclassified. Accuracy is simply the percentage of correctly classified instances over the total number of instances and is calculated as

$$ACC = \frac{(TP + TN)}{(TP + FP + TN + FN)}$$

where TP is the number of true positives, TN is the number of true negatives, FP is the number of false positives or over-predictions, and FN is the number of false negatives or missed predictions.

Accuracy is not an optimal measure of model performance if the data set is unbalanced (i.e. sizes of the classes are unequal) or if certain errors are to be considered more serious than others (e.g. false negatives compared to false positives).[Lapins et al., 2013] For this reason, all datasets were rebalanced to contain roughly equal numbers of active and inactive inhibitors.

As a further interrogation of results, the true positive rate and the true negative rate are also calculated for all models. The true positive rate is the fraction of true positives out of all actual inhibitors and it is calculated as follows

$$TruePositiveRate = \frac{TP}{(TP + FN)}$$

The true positive rate is also referred to as sensitivity and in MOE software it is called Accuracy on Active.

The true negative rate is the fraction of true negatives over the number of actual noninhibitors and is calculated as :

$$TrueNegativeRate = \frac{TN}{(TN + FP)}$$

The true negative rate is also known as the specificity of a classification, and in MOE it is called Accuracy on Inactives.

Figure 2.5: General Data Analysis Workflow

Binary QSAR after PLS of descriptors

The test sets were loaded and the washed structures were appended to the .csv files. All models were evaluated using the menu driven workflow in MOE and the classification probabilities were appended to the database file and saved as a .csv.

The resulting file was loaded into a MS Excel spreadsheet and the predictions classified as actives or inactives. Predicted probabilities of ≥ 0.5 were evaluated as active inhibitor predictions. Confusion matrices were then tabulated for predicted vs actual actives and inactives within the spreadsheet. And from the confusion matrices accuracy scores were calculated - total accuracy, specificity, and accuracy of inactive predictions. These results were saved and the results reported below.

κ -NN utilizing on the full set of 2-D descriptors

The full models follow a similar procedure to the previous workflow, except they omit data reduction by principle component analysis and use the full descriptor set to construct and evaluate models.

For each isozyme separately, training data is loaded into a dataframe. The Activity Score is assigned the role of response variable and all 186 are identified as predictor variables. Predictor variables are scaled and normed to mean 0 and standard deviation of 1. kNN.fit method is called from the

scikit-learn library to train a model, and then reports a the confusion matrix and an accuracy score for the training model by five-fold cross-validation.

Random Forest classification utilizing the full set of 2-D descriptors

For each isozyme separately, training data is loaded into a dataframe. The Activity Score is assigned the role of response variable and all 186 are identified as predictor variables. Predictor variables are scaled and normed to mean 0 and standard deviation of 1. RF.fit method is called from the scikit-learn library to train a model, and then reports a the confusion matrix and an accuracy score for the training model by five-fold cross-validation.

Support Vector Machine classification utilizing the full set of 2-D descriptors

For each isozyme separately, training data is loaded into a dataframe. The Activity Score is assigned the role of response variable and all 186 are identified as predictor variables. Predictor variables are scaled and normed to mean 0 and standard deviation of 1. SVM.fit method is called from the scikit-learn library to train a model, and then reports a the confusion matrix and an accuracy score for the training model by five-fold cross-validation.

A script was written for each isozyme that performed these three fit methods in series. Code and results are documented in Ipython notebooks that are currently hosted on github.com and freely accessible and downloadable. Presented in this way, they are easily verifiable and extendable. Experiment-

ing with other classification algorithms in scikit-learn simply requires adding new `method.fit` calls to the model building loop, because of the simple and consistent design of the scikit-learn library.

Chapter 3

Results

Results from MOE binary classification						
PCs		2c19	2c9	2d6	1a2	3a4
2	train	0.585	0.621	0.589	0.638	0.644
2	test	0.593	0.609	0.590	0.632	0.627
5	train	0.685	0.683	0.670	0.704	0.67
5	test	0.683	0.675	0.667	0.705	0.656
10	train	0.699	0.699	0.685	0.734	0.698
10	test	0.691	0.682	0.662	0.746	0.677
15	train	0.700	0.702	0.686	0.737	0.701
15	test	0.687	0.686	0.665	0.752	0.680
20	train	0.705	0.710	0.704	0.741	0.705
20	test	0.699	0.685	0.683	0.748	0.686
30	train	0.717	0.711	0.703	0.735	0.761
30	test	0.699	0.690	0.686	0.739	0.686
44	train	0.708	0.712	0.690	0.725	0.698
44pc	test	0.694	0.688	0.669	0.720	0.676

Table 3.1: Results from MOE binary classification

Modeling in Python

Toy Problem - κ NN on first 2 PCs

I don't know if I want to include this section. It is a gentle intro ending with a fairly easy to grasp visualization before the dimensionality explodes. So I'll probably include it.

After downloading the 2c19 dataset, the training data is loaded into memory. A script was written to perform a PCA on the training data, then select the first two principle components and construct a classification model using the κ -nearest neighbor algorithm. The major benefit of this exercise was the ease of visualization of results. Subsequently the two principle components used in model construction are plotted and overlayed with the model predictions.

The steps in this process script are as follows; the training data is loaded into a dataframe and Activity Score is identified as the response variables. All 186 molecular descriptors are selected as the input variables and then individually scaled to a mean of 0 and standard deviation of 1. PCA is performed on the scaled and normed descriptors. The first two PCs are used to train a κ -NN model that treats any ActivityScore over 40 as an active. The κ -NN function automatically performs a five-fold cross-validation of the model and creates a model object, which can be used to evaluate similarly formatted data, such as the test set.

Next the test set is loaded into a dataframe, the response variable identi-

fied, the descriptors scaled and normed, and subsequently evaluated by the model object. Accuracy in this case is calculated in the script, as well as a confusion matrix. These results are then plotted in the Ipython notebook with the principle components as axes, the predicted values from the model are calculated for every point and overlayed with the model results.

Here’s a scratch pad for all my results tables.

3.1 2C19 MOE models

2C19 Training Set Accuracy				2C19 Test Set Accuracy			
PC	Total	Active	Inactive	PC	Total	Active	Inactive
2	0.585	0.774	0.396	2	0.593	0.775	0.408
5	0.685	0.726	0.645	5	0.683	0.729	0.635
10	0.699	0.736	0.661	10	0.691	0.731	0.650
15	0.700	0.747	0.653	15	0.687	0.736	0.637
20	0.705	0.749	0.661	20	0.698	0.758	0.638
30	0.717	0.761	0.672	30	0.699	0.764	0.633
44	0.708	0.777	0.640	44	0.694	0.774	0.613

Table 3.2: 2C19 MOE Model Accuracy

Cytochrome P450 isozyme 2C19 models were trained on 4721 actives and 4741 inactives in the training set, and validated against 1193 actives and 1173 inactives in the test set. All models are considered good models when the total accuracy is above 0.6. Validating these models against the test set does not produce wildly divergent accuracy scores, indicating that overfitting is unlikely. The test set results represent compounds unseen by the model.

A visual scan of figure shows that model accuracy is higher for predicting




Figure 3.1: 2C19 MOE Model Accuracy

active compounds than inactive compounds, and total accuracy is somewhere in between. 20 principal components (PC20) appears to be an optimal number for capturing relevant variance while minimizing Type II errors. At 20 PCs, the total accuracy on the test set is 0.698. Accuracy on Actives is 0.758 and accuracy on inactives is 0.638.

3.2 2C9 MOE Models

Cytochrome P450 isozyme 2C9 models were trained on 3266 actives and 3327 inactives in the training set, and validated against 855 actives and 794 inactives in the test set. All models are considered good models when the total accuracy is above 0.6. Validating these models against the test set does not produce wildly divergent accuracy scores, indicating that overfitting is unlikely. The test set results represent compounds unseen by the model.

2C9 Training Set Accuracy				2C9 Test Set Accuracy			
PC	Total	Active	Inactive	PC	Total	Active	Inactive
2	0.620	0.695	0.547	2	0.609	0.702	0.510
5	0.683	0.722	0.645	5	0.675	0.713	0.634
10	0.699	0.744	0.655	10	0.682	0.730	0.632
15	0.702	0.752	0.653	15	0.686	0.739	0.630
20	0.710	0.756	0.665	20	0.685	0.731	0.636
30	0.711	0.754	0.668	30	0.690	0.730	0.646
44	0.712	0.770	0.655	44	0.688	0.750	0.621

Table 3.3: 2C9 MOE model results




Figure 3.2: 2C19 MOE Model Accuracy

A visual scan of figure shows that model accuracy is higher for predicting active compounds than inactive compounds, and total accuracy is somewhere in between. 20 principal components (PC20) appears to be an optimal number for capturing relevant variance while minimizing Type II errors. At 20PCs, the total accuracy on the test set is 0.685. Accuracy on Actives is 0.731 and accuracy on inactives is 0.636.

3.3 1A2 MOE Models

1A2 Training Set Accuracy				1A2 Test Set Accuracy			
PC	Total	Active	Inactive	PC	Total	Active	Inactive
2	0.638	0.752	0.522	2	0.632	0.758	0.517
5	0.704	0.746	0.662	5	0.705	0.773	0.643
10	0.734	0.755	0.712	10	0.746	0.749	0.743
15	0.737	0.770	0.704	15	0.752	0.775	0.731
20	0.741	0.781	0.701	20	0.748	0.779	0.721
30	0.735	0.782	0.686	30	0.739	0.783	0.698
44	0.725	0.777	0.672	44	0.720	0.773	0.672

Table 3.4: 1A2 MOE model results

Cytochrome P450 isozyme 1A2 models were trained on 4820 actives and 4780 inactives in the training set, and validated against 529 actives and 580 inactives in the test set. All models are considered good models when the total accuracy is above 0.6. Validating these models against the test set does not produce wildly divergent accuracy scores, indicating that overfitting is unlikely. The test set results represent compounds unseen by the model.

A visual scan of figure shows that model accuracy is higher for predicting active compounds than inactive compounds, and total accuracy is somewhere in between. 20 principal components (PC20) appears to be an optimal number for capturing relevant variance while minimizing Type II errors. At 20PCs, the total accuracy on the test set is 0.748. Accuracy on Actives is 0.779 and accuracy on inactives is 0.721.




Figure 3.3: 1A2 MOE Model Accuracy

3.4 2D6 MOE Models

2D6 Training Set Accuracy				2D6 Test Set Accuracy			
PC	Total	Active	Inactive	PC	Total	Active	Inactive
2	0.589	0.753	0.426	2	0.590	0.738	0.443
5	0.670	0.706	0.634	5	0.667	0.685	0.650
10	0.685	0.721	0.648	10	0.662	0.688	0.636
15	0.686	0.709	0.662	15	0.665	0.692	0.639
20	0.705	0.721	0.689	20	0.683	0.699	0.668
30	0.703	0.732	0.673	30	0.686	0.696	0.677
44	0.690	0.725	0.656	44	0.669	0.707	0.632

Table 3.5: 2D6 MOE model results

Cytochrome P450 isozyme 2D6 models were trained on 2219 of actives and 2214 of inactives in the training set, and validated against 552 actives and 557 inactives in the test set. All models are considered good models when the total accuracy is above 0.6. Validating these models against the test set does not produce wildly divergent accuracy scores, indicating that

overfitting is unlikely. The test set results represent compounds unseen by the model.




Figure 3.4: 2D6 MOE Model Accuracy

A visual scan of figure shows that model accuracy is higher for predicting active compounds than inactive compounds, and total accuracy is somewhere in between. 20 principal components (PC20) appears to be an optimal number for capturing relevant variance while minimizing Type II errors. At 20 PCs, the total accuracy on the test set is 0.683. Accuracy on Actives is 0.699 and accuracy on inactives is 0.668.

3.5 3A4 MOE Models

Cytochrome P450 isozyme 3A4 models were trained on 4234 of actives and 4193 of inactives in the training set, and validated against 1033 actives and 1074 inactives in the test set. All models are considered good models when

3A4 Training Set Accuracy				3A4 Test Set Accuracy			
PC	Total	Active	Inactive	PC	Total	Active	Inactive
2	0.644	0.699	0.589	2	0.627	0.683	0.573
5	0.675	0.727	0.623	5	0.656	0.719	0.595
10	0.698	0.741	0.655	10	0.677	0.719	0.637
15	0.701	0.746	0.656	15	0.680	0.733	0.628
20	0.705	0.749	0.660	20	0.686	0.736	0.638
30	0.709	0.761	0.657	30	0.686	0.742	0.631
44	0.698	0.773	0.623	44	0.676	0.747	0.607

Table 3.6: 3A4 MOE model results

the total accuracy is above 0.6. Validating these models against the test set does not produce wildly divergent accuracy scores, indicating that overfitting is unlikely. The test set results represent compounds unseen by the model.




Figure 3.5: 3A4 MOE Model Accuracy

A visual scan of figure shows that model accuracy is higher for predicting active compounds than inactive compounds, and total accuracy is somewhere in between. 20 principal components (PC20) appears to be an optimal number

for capturing relevant variance while minimizing Type II errors. At 20PCs, the total accuracy on the test set is 0.686. Accuracy on Actives is 0.736 and accuracy on inactives is 0.638.

3.6 Comparison of Methods

3.6.1 1A2

1A2 Classification Method Comparison		
	Training Set	Test Set
MOE 20PC	0.741	0.748
kNN	0.761	0.754
Random Forest	0.770	0.769
SVD	0.800	0.804

Table 3.7: Comparison of Classification Methods for 1A2

3.6.2 2C9

2C9 Classification Method Comparison		
	Training Set	Test Set
MOE 20PC	0.710	0.685
kNN	0.721	0.692
Random Forest	0.717	0.687
SVD	0.749	0.720

Table 3.8: Comparison of Classification Methods for 2C9

2C19 Classification Method Comparison		
	Training Set	Test Set
MOE 20PC	0.705	0.698
kNN	0.730	0.720
Random Forest	0.736	0.721
SVD	0.767	0.756

Table 3.9: Comparison of Classification Methods for 2C19

3.6.3 2C19

3.6.4 2D6

2D6 Classification Method Comparison		
	Training Set	Test Set
MOE 20PC	0.705	0.683
kNN	0.717	0.678
Random Forest	0.729	0.707
SVD	0.755	0.725

Table 3.10: Comparison of Classification Methods for 2D6

3.6.5 3A4

3A4 Classification Method Comparison		
	Training Set	Test Set
MOE 20PC	0.705	0.686
kNN	0.716	0.674
Random Forest	0.719	0.683
SVD	0.755	0.714

Table 3.11: Comparison of Classification Methods for 3A4




Figure 3.6: Inhibitor Classification Method Comparison on the Training Set




Figure 3.7: Inhibitor Classification Method Comparison on the Test Set

Chapter 4

Discussion

4.1 Other attempts at modelling Assay 1851

Pubchem Bioassay AID1851 has been the basis for several attempts to advance *in silico* screening development. Several research groups have developed large-scale single target models for the single isoforms, [Vasanthanathan et al., 2009, Sridhar et al., 2012, Sun, 2012, Novotarskyi et al., 2010] and Cheng et al created single target QSAR models for all five CYP isoforms which are the most comparable to my efforts. [Cheng et al., 2011] And Lapins et al. also pursued all five isozymes by a proteochemometric method that also takes into account each isozyme's protein sequence. [Lapins et al., 2013] These models all showed good predictive performances according to our standards of a 0.6 accuracy threshold.

4.1.1 Cheng’s results compared to mine

It is highly desirable to develop computational models that can predict the inhibitive effect of a compound against a specific CYP isoform. In this Cheng’s study inhibitor predicting models were developed for five major CYP isoforms, namely 1A2, 2C9, 2C19, 2D6, and 3A4, using a combined classifier algorithm on a large data set containing more than 24,700 unique compounds, extracted from PubChem. The combined classifiers algorithm is an ensemble of different independent machine learning classifiers including support vector machine, C 4.5 decision tree, */kappa*-nearest neighbor, and naive Bayes, joined by a back-propagation artificial neural network (BP-ANN). Those models were validated by 5-fold cross-validation with a diverse validation set composed of about 9000 diverse unique compounds. The range of the AUROC for the validation set was 0.764 to 0.886. The overall performance of combined classifiers fused by BP-ANN was superior to that of three classic fusion techniques (Mean, Maximum, and Multiply). Cheng et al. claim these classification models are applicable for virtual screening of the five major CYP isoforms inhibitors or can be used as simple filters of potential chemicals in drug discovery.[Cheng et al., 2011]

4.2 Sources of error

This assay required recombinant sources of CYP enzymes, because the probes are not sufficiently P450 isoform-selective to be used with those derived

from human liver microsomes. Another potential source for false positives in these assays can be compounds that interfere with light generation directly, such as compounds that interfere with luciferase enzymatic activity. [Zlokarnik et al., 2005]

The results of the large-scale screening of Pubchem AID1851 against five CYP isoforms identified that the majority of compounds in a typical chemical library cross-inhibited several isoforms, while only a small fraction of the compounds did not inhibit any of the isoforms. [Veith et al., 2009]

The phenomenon of uniformly lower accuracy on the inactives versus the actives might be explained by the composition of the inactive class. Both inactive inhibitors as well as compounds of uncertain activity were both included under the label 'inactive' for binary classification purposes. The chosen representation of class inclusion may have less potential for successful discrimination than. To test assumptions about class labels, the number of classes could be increased or the Activity Score threshold could be altered in future models.

4.3 Machine Learning

Developing successful machine learning applications still requires a substantial amount of black art that is hard to find in textbooks.[Domingos, 2012]

Unlike in this study, raw data is often not in a form that is amenable to learning, but you can construct features from it that are. Easily the

most important factor in a good model is the features used. If there are many independent features that correlate well with the class, learning is easy. However if the class is a very complex function of the features, the time required to train a good model may take too long or be impossible. [Domingos, 2012] Most of the time and effort in a machine learning project goes typically into feature engineering, and it is one area where domain expertise and experience with chemical library design can mean the difference between success and failure.

Also machine learning is not a one-shot process of building a data set and running a learner, but rather an iterative process of running the learner, analyzing results, modifying the data and/or learner, and repeating. Learning is often the quickest part of this, but that's because a lot of effort has gone into sharing and codifying that knowledge. Feature engineering is more difficult because it is domain-specific, while learners can be largely general-purpose. [Domingos, 2012]

There is ultimately no replacement for the smarts you put into feature engineering. On the other hand, running a learner with a very large number of features to find out which ones are useful in combination may be too time-consuming, or cause overfitting. [Domingos, 2012]

As a rule, it pays to try the simplest learners first (e.g. naive Bayes before logistic regression, κ -nearest neighbor before support vector machines). More sophisticated learners may be seductive, but they are usually harder to use, because there are more tuning parameters and their internals are more

opaque.[Domingos, 2012]

One way to divide learning algorithms is as follows: those whose representation has a fixed size, like linear classifiers, and those whose representation can grow with the data, like decision trees. Fixed-size learners can only take advantage of so much data. Variable-sized learners can in principle learn any function given sufficient data, but in practice may not, because of limitations of the algorithm or computational cost. Also, because of the curse of dimensionality, no existing amount of data may be enough.[Domingos, 2012] Each type of learner comes with its own assumptions and, as of yet, none are definitely best at all situations, but some are clearly better than most. [Hand, 2006, Fernández-Delgado et al., 2014]

Perhaps because each type of learner is looking at the data from different angles, combining the results of different learners into an ensemble of models is a technique that has demonstrated its effectiveness at improving overall accuracy. Creating model ensembles is now standard in machine learning. The statistic community uses techniques, such as bagging, boosting and stacking to resample training data or reweight classifier inputs on previously misclassified data in order to reduce variance while minimizing bias. But in the simplest form of an ensemble of models can be thought of as a 'majority vote', where final class assignment rest at which most of the included models arrived.

Ensembles of models frequently outperform any single model, so a logical next step of this study would be to combine results from all models developed.

Chapter 5

Conclusion

QSAR models and their value in informing regulatory decision making will likely increase with the standardization of analytical approaches, more complete and reliable data collection methods and a better understanding of toxicity mechanisms in the role of disease as well as individual susceptibility to adverse clinical events. [Kruhlak et al., 2012]

I have demonstrated several methods for building models for binary classification for QSAR of cytochrome P450 isozymes. They all do pretty well and are significantly better than randomly guessing. There are many more machine learning methods to try, of particular interest are neural networks and more Bayesian methods. The data and code for model comparison for this study is available online, so there are no artificial barriers to benchmarking further results.

Bibliography

- [Berg, 2014] Berg, E. L. (2014). Systems biology in drug discovery and development. *Drug discovery today*, 19(2):113–25.
- [Boulesteix and Schmid, 2014] Boulesteix, A. L. and Schmid, M. (2014). Machine learning versus statistical modeling. *Biometrical Journal*, 4(56):588–593.
- [Breiman, 2001] Breiman, L. (2001). Statistical Modeling : The Two Cultures. 16(3):199–231.
- [Cheng et al., 2011] Cheng, F., Yu, Y., Shen, J., Yang, L., Li, W., Liu, G., Lee, P. W., and Tang, Y. (2011). Classification of cytochrome P450 inhibitors and noninhibitors using combined classifiers. *Journal of chemical information and modeling*, 51(5):996–1011.
- [Cohen, 2004] Cohen, J. E. (2004). Mathematics is biology’s next microscope, only better. *PLoS Biol*, 2(12):e439.
- [Danielson, 2002] Danielson, P. B. (2002). The cytochrome P450 superfamily: biochemistry, evolution and drug metabolism in humans. *Current drug metabolism*, 3(6):561–97.
- [Domingos, 2012] Domingos, P. (2012). A few useful things to know about machine learning. *Communications of the ACM*, 55(10):78.
- [Fernández-Delgado et al., 2014] Fernández-Delgado, M., Cernadas, E., Barro, S., and Amorim, D. (2014). Do we need hundreds of classifiers to solve real world classification problems? *Journal of Machine Learning Research*, 15:3133–3181.
- [Hand, 2006] Hand, D. J. (2006). Classifier technology and the illusion of progress. *Statist. Sci.*, 21(1):1–14.

- [Hansch and Fujita, 1964] Hansch, C. and Fujita, T. (1964). Analysis. A Method for the Correlation of Biological Activity and Chemical Structure. *Journal of American Chemical Society*, 86:1616–1626.
- [James et al., 2013] James, G., Witten, D., Hastie, T., and Tibshirani, R. (2013). *An Introduction to Statistical Learning*. Springer.
- [Kruhlak et al., 2012] Kruhlak, N. L., Benz, R. D., Zhou, H., and Colatsky, T. J. (2012). (Q)SAR modeling and safety assessment in regulatory review. *Clinical pharmacology and therapeutics*, 91:529–34.
- [Labute, 1999] Labute, P. (1999). Binary qsar: A new method for determination of quantitative structure activity relationships. *Pacific Symposium on Biocomputing*, 4:444–455.
- [Lapins et al., 2013] Lapins, M., Worachartcheewan, A., Spjuth, O., Georgiev, V., Prachayasittikul, V., Nantasenamat, C., and Wikberg, J. E. S. (2013). A unified proteochemometric model for prediction of inhibition of cytochrome p450 isoforms. *PloS one*, 8(6):e66566.
- [Nantasenamat, 2009] Nantasenamat, C. (2009). A practical overview of quantitative structure-activity relationship. *EXCLI J*, 8:74–88.
- [Novotarskyi et al., 2010] Novotarskyi, S., Sushko, I., Körner, R., Pandey, A., and Tetko, I. (2010). Classification of CYP450 1A2 inhibitors using PubChem data. *Journal of Cheminformatics*, 2(Suppl 1):P40.
- [Prli and Lapp, 2012] Prli, A. and Lapp, H. (2012). The plos computational biology software section. *PLoS Comput Biol*, 8(11):e1002799.
- [Singh et al., 2011] Singh, D., Kashyap, A., Pandey, R. V., and Saini, K. S. (2011). Novel advances in cytochrome P450 research. *Drug discovery today*, 16(17-18):793–9.
- [Sridhar et al., 2012] Sridhar, J., Liu, J., Foroozesh, M., and Stevens, C. L. K. (2012). Insights on cytochrome p450 enzymes and inhibitors obtained through QSAR studies. *Molecules (Basel, Switzerland)*, 17(8):9283–305.
- [Sun, 2012] Sun, H. (2012). Predictive Models for Cytochrome P450 Isozymes Based on Quantitative High Throughput Screening Data. 51(10):2474–2481.

- [Vasanthanathan et al., 2009] Vasanthanathan, P., Taboureau, O., Oostenbrink, C., Vermeulen, N. P. E., Olsen, L., and Jørgensen, F. S. (2009). Classification of Cytochrome P450 1A2 Inhibitors and Noninhibitors by Machine Learning Techniques ABSTRACT :. 37(3):658–664.
- [Veith et al., 2009] Veith, H., Southall, N., Huang, R., and James, T. (2009). Comprehensive characterization of cytochrome P450 isozyme selectivity across chemical libraries. *Nature ...*, 27(11):1050–1055.
- [Visser et al., 2014] Visser, S. a. G., de Alwis, D. P., Kerbusch, T., Stone, J. a., and Allerheiligen, S. R. B. (2014). Implementation of quantitative and systems pharmacology in large pharma. *CPT: pharmacometrics & systems pharmacology*, 3(October):e142.
- [Zlokarnik et al., 2005] Zlokarnik, G., Grootenhuis, P. D. J., and Watson, J. B. (2005). High throughput P450 inhibition screens in early drug discovery Case study REVIEWS. 10(21):1443–1450.