

Cytochrome P450 Inhibitor Classification with Statistical Learning

by
Daniel Siegle

A Thesis Submitted to
The Faculty of the Graduate School at
North Carolina Central University
In Partial Fulfillment of the Requirements
for the Degree of Master of Science

Durham
2015

Approved by

Committee Chair: Weifan Zheng

Committee Member: Tomas Ding

Committee Member: ClarLynda Williams-DeVane

Abstract

SIEGLE, DANIEL E, M.S. "Cytochrome P450 Inhibitor Classification with Statistical Learning (2015)" Directed by Dr. Weifan Zheng, 53pp.

This project compares methods of Cytochrome P450 inhibitor prediction based on compound structures. CYPs are a natural first choice in which to develop *in silico* models because of their central role in drug candidate rejection due to adverse drug-drug interactions. Several non-linear, high-dimensional classification models were built and compared using a large, publicly available high throughput screening luminescence assay (PubChem AID1851) against five CYP isozymes (1A2, 2C9, 2C19, 2D6 and 3A4). The methods compared are a Bayesian binary QSAR method from the Molecular Operating Environment, and 3 standard machine learning methods implemented in the Python programming language; κ -Nearest Neighbor, Random Forests, and Support Vector Machines. They all performed well, with the methods implemented in freely available software performing as well or better than the method in software that is widely accepted in industry.

Contents

List of Tables	iii
List of Figures	iv
1 Background	1
1.1 Drug Discovery Productivity Challenges	1
1.2 Cytochrome P450 Superfamily	3
1.3 Early Compound Profiling <i>In Silico</i>	7
1.4 Quantitative Structure-Activity Relationship	9
1.5 Statistical Machine Learning	13
1.6 Sources of Data for Learning	14
1.7 Reproducibility	15
1.8 Aims	17
2 Materials and Methods	18
2.1 Review of PubChem Assay 1851	18
2.2 Data Set Preparation and Molecular Descriptor Generation	20
2.3 Feature selection	23
2.4 Modeling - Mapping Descriptors to Activity Data	24
2.4.1 Binary QSAR in the Molecular Operating Environment	24
2.4.2 κ -Nearest Neighbor (κ NN)	26
2.4.3 Random Forest (RF)	26
2.4.4 Support Vector Machines (SVM)	27
2.5 Model Evaluation and Validation	27
3 Results	34
3.1 1A2 MOE Models	34
3.2 2C9 MOE Models	35
3.3 2C19 MOE models	37
3.4 2D6 MOE Models	38
3.5 3A4 MOE Models	39
3.6 1A2 Classification Method Comparison	41
3.7 2C9 Classification Method Comparison	41
3.8 2C19 Classification Method Comparison	42
3.9 2D6 Classification Method Comparison	42

3.10	3A4 Classification Method Comparison	43
3.11	Overall Method Comparison	44
4	Discussion	46
4.1	Other Attempts at Modeling Assay 1851	46
4.2	Sources of Error	47
4.3	Machine Learning for Pharmaceutical Sciences	48
5	Conclusion	50

List of Tables

3.1	1A2 MOE Model Results	34
3.2	2C9 MOE Model Results	35
3.3	2C19 MOE Model Results	37
3.4	2D6 MOE Model Results	38
3.5	3A4 MOE Model Results	39
3.6	Comparison of Classification Methods for 1A2	41
3.7	Comparison of Classification Methods for 2C9	41
3.8	Comparison of Classification Methods for 2C19	42
3.9	Comparison of Classification Methods for 2D6	42
3.10	Comparison of Classification Methods for 3A4	43

List of Figures

1.1	Ribbon Diagram of CYP1A2, PDB 2HI4	4
1.2	Mechanism of Action of Cytochrome P450 Catalytic Heme Group . .	5
1.3	Ribbon Diagram of CYP3A4 with Heme Center	6
1.4	Typical QSAR Workflow	12
2.1	Data Download and Preparation	20
2.2	Isozyme Data Set Preparation	22
2.3	Supervised Classification Model Training and Validation	24
2.4	Supervised Learning Workflow for Clasification	33
3.1	1A2 MOE Model Accuracy	35
3.2	2C19 MOE Model Accuracy	36
3.3	2C19 MOE Model Accuracy	37
3.4	2D6 MOE Model Accuracy	39
3.5	3A4 MOE Model Accuracy	40
3.6	Inhibitor Classification Method Comparison on the Training Set . . .	44
3.7	Inhibitor Classification Method Comparison on the Test Set	45

Acknowledgements

To my advisor and committee chair Dr. Weifan Zheng for all his guidance and support. Thanks also to committee members Dr. ClarLynda Williams-DeVane and Dr. Tomas Ding for their contributions and consideration.

I would also like to thank Dr. Johnathan Sexton for advising my first foray into research and special thanks to Natacha Janvier-Derilus for unflagging professionalism in advising and dedication to the students of the BRITE program.

Chapter 1

Background

1.1 Drug Discovery Productivity Challenges

The goal of pharmaceutical sciences is to identify safe and efficacious drugs for the market. Toxicity is a large contributor to drug candidate attrition in drug development. Inhibition of enzymes in the cytochrome P450 superfamily is a major source of toxicity in human and animal models because of their role in first pass metabolism. If a compound inhibits cytochrome P450 it is likely to lead to toxic effects when first pass metabolism fails to clear other therapeutic compounds or alters their pharmacokinetics. Because the discovery and development of a each new pharmaceutical drug requires more than a billion dollars across an average of 12 years, it would be useful to know whether a compound of interest inhibits cytochrome P450s as early as possible.

The Cytochrome P450 superfamily is large and varied. Different isozymes metabolize different substrates. Even within and between individuals pharmacokinetic and pharmacodynamic variability can be high for reasons not entirely characterized. Designing assays for every isozyme and SNP variant and running them against every compound of interest as a general strategy is cost prohibitive. This project was de-

signed to test the predictive power of computational methods for cytochrome P450 inhibition potential based simply on knowledge of chemical structure.

There has been a downward trend in drug development productivity for the last three decades for a variety of reasons. The difficulty of improving upon current treatments in novel ways, increasingly cautious regulatory environments, and over-reliance on early stage identification of 'silver bullet' treatments that end up as late stage failures have all been suggested as causes of this productivity decline. [Scannell et al., 2012] And integrated computational approaches have been proposed as one way to control costs in ongoing pharmaceutical research and development. [Visser et al., 2014] The foundational skills demonstrated in this thesis are needed to pursue a systems approach to drug development that industry has recently turned toward as a way to boost R&D productivity. [Berg, 2014] This thesis is also part of the larger project of *in silico* drug-development, which is attempting to reduce reliance on exploratory *in vivo* and clinical drug testing and increase the number of effective treatments for patients while decreasing the amount of time to develop them.

The Quantitative Structure Activity Relationship (QSAR) approach to associating compound structure with bioactivity that goes back at least to Hansch [Hansch and Fujita, 1964], has a long history with *in silico* drug-development efforts. QSAR started with direct measures of chemical compounds and then later derived features, and used them to then build expert systems or statistical models that tried to predict biological activity. During those same decades, the field of machine learning emerged; using computation in an analogous and more general way – to associate features with results, inputs with outputs. Machine learning can be thought of as using algorithms to figure out how to perform important tasks by generalizing from examples. These algorithms are of course usually laborious, which necessitates their execution by computer.

Statistical machine learning builds upon the peer prediction of machine learning.

It also allows for prediction but focuses more on models and methods that can be used by scientists and engineers. [James et al., 2013] Further extension of statistical learning to the pharmaceutical sciences can lead to important contributions to systems pharmacology. Or rather, statistical learning methods are an important precursor to the needs of systems biology and systems pharmacology modeling.

This project compares different techniques for cytochrome P450 inhibition prediction in the framework of statistical learning. In order to demonstrate generality and applicability to the pharmaceutical sciences, the subjects of this study are the five isozymes of Cytochrome P450 (CYP) that are involved in metabolism of 90% of all therapeutic drugs (CYPs 1A2, 2C9, 2C19, 2D6 and 3A4).

Because replicability is one of the main principles of the scientific method, as much as possible, the code and data used in this study is source controlled and publicly available. Reproducibility as a practice is a habit, and a good one to get into. Every attempt was made to make the materials and methods for this project reproducible in a completely automated way to allow for validation of results and extension of the work.

1.2 Cytochrome P450 Superfamily

One of the largest and most functionally diverse protein superfamilies is the cytochrome P450 family of hemoproteins. From bacteria to humans, the functional breadth of cytochrome P450 activity is far ranging. At the latest count there were significantly more than 2000 identified cytochrome P450 genomic and cDNA sequences that have been divided into a total of 265 different families. [Danielson, 2002] Cytochromes P450 appear in every kingdom from bacteria to higher eukaryotes. Multiple cytochrome P450 genes can be expressed simultaneously as different isozymes and the number of genes per species is highly variable with a tendency for higher eukaryotes

to possess more. The cytochromes P450 constitute the major enzyme family capable of catalyzing the oxidative biotransformation of most drugs and other lipophilic xenobiotics and are therefore of particular relevance for clinical pharmacology. The central role that these ubiquitous proteins play as phase I enzymes in human drug metabolism makes them very important to the pharmaceutical industry.




Figure 1.1: Ribbon Diagram of CYP1A2, PDB 2HI4

CYP families are classified based on pairwise amino acid sequence identity between individual members. Families CYP 1-3 are involved in phase I metabolism of human drugs and xenobiotic compounds, whereas other CYP families (CYP 4, 11, 17, and 21) are involved in the metabolism of endogenous compounds such as fatty acids, steroids, eicosanoids, bile acids and fat soluble vitamins. [Singh et al., 2011]




Figure 1.2: Mechanism of Action of Cytochrome P450 Catalytic Heme Group

Figure 1.2 illustrates the catalytic cycle that occurs at the central heme group in all Cytochromes P450 as follows:

1. First the substrate binds and induces a conformational change in the enzyme that usually displaces a water group. This brings the substrate in close proximity to the heme group.
2. Bound substrate induces electron transfer from a nearby reductase that provides NAD(P)H. At this point the -CO complexed with the heme group absorbs light spectra at max 450nm, noted by the 450 in the original naming of Cytochromes P450.
3. A second oxygen molecule displaces a carbon followed shortly by

4. A second electron transferred from a nearby reductase that leads to
5. Rapid double protonation of the peroxo group that results in one molecule of water being released, leaving the ferrous center with one double bound oxygen molecule. An alternative route this chemical state is shown by the S step in Figure 1.2, called the 'peroxide shunt', that entails oxidation of the heme center directly.
6. At this point, it depends on the substrate and isozyme involved as to what sort of reaction is catalyzed, but the figure shows a hydroxylation, which is a very common way to make lipophilic substances more water soluble.




Figure 1.3: Ribbon Diagram of CYP3A4 with Heme Center Highlighted, PDB 1W0E

There is a considerable substrate overlap between enzymes of this superfamily.

Being broadly specific with respect to their substrates, CYPs are therefore susceptible to inhibition by a large variety of chemical compounds. It follows that the CYP enzymes that are involved in the oxidative metabolism of drugs play a major part in the activation and elimination of therapeutic drug molecules. CYP inhibition that leads to decreased elimination and/or changed metabolic pathways of other substrates is a major cause of adverse drug-drug interactions. [Lapins et al., 2013] So adverse side effects of drug-drug interactions are an important area of inquiry, especially during the research phase of drug discovery.

1.3 Early Compound Profiling *In Silico*

Drug discovery is a multi-parameter optimization process in which compounds are designed for interaction with their target while simultaneously minimizing off-target activities. [Zlokarnik et al., 2005] In the normal course of taking a compound from hit to lead to drug candidate, adding drug-like properties may minimize the risk of making potent and target-class selective compounds that are biologically inaccessible, but at the same time does little to address the combinatorial complexities of specific drug-drug interactions due to off-target effects.

Because of the complex nature of toxicity, safety prediction is considered more challenging than efficacy prediction for a number of reasons. Toxicity mechanisms may be unknown or poorly characterized in higher organisms, and similar pathways and targets may be associated with different toxicities and adverse events. Toxicity prediction must also encompass a number of complex interactions and remain alert to the possibility of finding the unexpected. For instance, toxicities could result from on-target effects due to incomplete knowledge or inadequate target validation, or from off-target effects mediated via unknown molecules and mechanisms, or even from genetic variation or comorbidities in any of the previously mentioned pathways.

[Kruhlak et al., 2012]

Techniques for high-throughput *in vitro* screening of CYP inhibition have been developed and implemented broadly in drug discovery pipelines across pharmaceutical companies and research institutions, resulting in the generation of large sets of data. Some of this accumulation of data has been released through academic research initiatives (e.g. PubChem Bioassays AID 410 and 1851). These collections enable development of structure-activity relationship models for *in silico* prediction of CYP inhibition by a much larger pool of researchers than those who designed and carried out the assays.

The promise of *in silico* screening that remains very appealing is that, with a steady increase in computing power, screening costs could become negligible. The hope is that virtual compounds could be screened for CYP liabilities in order to realize savings and reduce the number of candidates with questionable prospects that would otherwise have to be synthesized. [Zlokarnik et al., 2005] Achieving this would also turn the costly animal toxicity test phase of preclinical research into a validation step rather than a screening step.

In order to generate the necessary data, many high-throughput technologies are now available to detect P450 inhibitors. High-throughput screening data can be used to guide medicinal chemists away from these interactions at an early stage. In certain cases it might also identify the inhibition issue allowing intervention by targeted modification of the CYP interacting functionality. According to Zlokarnik, to be generally useful, P450 inhibition screens need to be calibrated against standard methods and preferably also tested with a large set of drugs, for which human drug-drug interaction outcome is known. [Zlokarnik et al., 2005] This should also decrease the number of withdrawals of novel drugs from the market due to inhibition of major P450 isozymes.

The ability to predict clinical safety based on chemical structures has recently

become an increasingly important part of regulatory decision-making. QSAR models are currently used by industry and by regulators to evaluate safety concerns and possible nonclinical effects of a drug when adequate safety data is absent or inconclusive. [Kruhlak et al., 2012]

As an example, the United States of America is about to become the last country to accept QSAR in the drug approval process. The drafting of International Committee on Harmonization (ICH) M7 guideline can be viewed as setting a precedent for possible future, broader regulatory applications of QSAR modeling. ICH M7, will for the first time specify that – under very specific conditions – the results of QSAR computational toxicology predictions will be considered sufficient for genotoxic contaminants of pharmaceuticals under consideration and thereby eliminate the need for laboratory testing. [Kruhlak et al., 2012]

1.4 Quantitative Structure-Activity Relationship

Quantitative structure-activity relationship modeling is generally accepted as the construction of predictive models of biological activities as a function of structural and molecular information of a compound or compound library. [Nantasenamat, 2009]

QSAR models describe the correlation between molecular features and activity at a given end point of interest. There have also been attempts to make structure activity relationship (SAR) models constructed by using human expert knowledge (expert rule-based), but QSAR models are typically defined as those that use mathematical methods to analyze the statistical correlations between molecular features and activity.

A QSAR model that defines the mathematical relationships between descriptors and biological activities of known molecules, differs from receptor binding-based efficacy prediction which takes into account binding site characteristics as well as molec-

ular docking analysis. In contrast to QSAR, receptor-binding methods attempt to predict drug efficacy based on known mechanisms of action and medicinal chemistry by individually studying molecular interactions between a drug and targets/receptors. [Kruhlak et al., 2012]

It follows from the 'similarity principle' that new and untested compounds possessing similar molecular features as known compounds are assumed to possess similar activities and properties. In this way, QSAR models can make it possible to predict the biological activities of a given compound as a function of its molecular structure. Several successful models have been published over the years which encompass a wide spectrum of biological and physicochemical properties.

Applied QSAR, as described above, has typically been used for drug discovery and development but has also been used to correlate molecular information with other physiochemical properties. This later approach is termed quantitative structure-property relationship (QSPR). Derived molecular parameters can account for hydrophobicity, topology, electronic properties, and steric effects among other things. These characteristics of compounds can either be determined empirically through experimentation or theoretically via computational chemistry as needed. [Nantasenamat, 2009] The parameters derived from compound structure are referred to as molecular descriptors.

Molecular Descriptors

Molecular descriptors can be thought of as the mathematical representation of essential information of a molecule in terms of its own physiochemical properties. Depending on the needs of the analysis, properties considered can be electronic, geometric, hydrophobic, constitutional, lipophilic, steric, solubility, quantum chemical, or topological. From a practical viewpoint, molecular descriptors are chemical information that is encoded within the molecular structures. [Nantasenamat, 2009]

Molecular features can be either experimentally measured or calculated values. They come in the form of simple physiochemical properties such as logP or logarithmic acid dissociation constant (pKa), numerical representations of substructure fragments, or purely mathematical descriptors. Mathematical descriptors are chemical structural features represented in numerical form, and range from simple atom counts to the product of complex equations that describe electron distribution across a molecule. [Kruhlak et al., 2012]

Molecular descriptors as predictors in QSAR modeling are typically less precise than the lock and key relationships that underpin the docking approach to computer-aided drug design. The basic assumption in QSAR modeling is that similar molecules exhibit similar biological activity so that physiochemical properties and/or structural properties of a molecule encoded as molecular descriptors can then be used to predict the biological activity of structurally related compounds with some degree of confidence.

Modeling

QSAR models can be described as global or local. Global models incorporate chemicals with a range of molecular features acting across the spectrum of chemical pathways, whereas local models are highly focused on a single chemical class and end point. Although local models generally have much higher accuracy, their narrow domain of applicability renders them impractical in most regulatory environments where predictions need to be made across a variety of molecules including active pharmaceutical ingredients, as well as metabolites, reagents, and synthetic intermediates. [Kruhlak et al., 2012]

The construction of QSAR models typically follows two main steps:

- Description of molecular structure with derivation of descriptors
- Multivariate analysis correlating molecular descriptors with observed activities.

Additional intermediate steps are also crucial for successful development of such QSAR models and include data preprocessing and statistical evaluation. See Figure 1.4 [Nantasenamat, 2009]




Figure 1.4: Typical QSAR Workflow [Nantasenamat, 2009]

A typical QSAR workflow treats chemical structure management, descriptor calculation, and statistical analyses as separate steps that are often performed by non-integrated software packages. This can lead to low throughput and even the lack of possibility of performing predictions for new compounds or the inability to update the models when new data become available, depending on the workflow. Approaches that integrate as many of these steps as possible are generally preferred.

According to Kruhlak, et al., the successful development of a QSAR model for safety prediction requires a sufficient amount of high-quality data, the appropriate selection of descriptors, the availability of one or more suitable statistical or mathematical models and an effective training and validation strategy. [Kruhlak et al., 2012]

1.5 Statistical Machine Learning

Machine learning algorithms figure out how to perform important tasks by generalizing from examples. A concise definition of statistics is as the applied science that constructs and studies techniques for data analysis. (Jan de Leeuw) Statistical learning refers to a set of approaches for estimating a function that describes a dataset as a precursor for prediction or inference. [James et al., 2013] Statistical machine learning constructs that function by generalizing from examples, i.e. data.

Leo Breiman wrote a landmark paper that documented the beginnings of this approach. He said

'There are two cultures in the use of statistical modeling to reach conclusions from data. One assumes that the data are generated by a given stochastic data model. The other uses algorithmic models and treats the data mechanisms as unknown' [Breiman, 2001]

He goes on to claim that the statistical community had traditionally preferred the first view.

Classification is a well understood area of machine learning. A classifier takes a system of inputs, typically a vector of discrete and/or continuous feature values and then outputs a single discrete value, the class. [Domingos, 2012] It can be used when there are multiple examples of items of interest that are then used to guide the determination of a new item based on its characteristics.

The current state of machine learning is fundamentally a subset of optimization and has found its biggest successes in fields where there are far more variables than parameters. The ultimate goal of machine learning in these cases is to generalize beyond examples in a training set, because no matter how many example the training set contains, it is unlikely that those exact examples will be seen again in practical applications. [Domingos, 2012]

In the context of QSAR then, with enough prior knowledge and sound assay results, machine learning may be a practical approach to fill in the gaps of clinical knowledge for any relevant CYP isozyme when queried against any untested compound of interest so long as the compound structure is known.

This is a wider view than traditionally encountered in most applied science education where simple hypothesis testing predominates. Opening up the field of data analysis like this brings opportunity for exploration but also new concerns, such as the 'curse of dimensionality', 'degrees of freedom of the analyst', 'black box algorithms', bias estimation difficulties and the 'no free lunch theorem'. [Boulesteix and Schmid, 2014]

1.6 Sources of Data for Learning

Systems biology research encompasses the generation of high-throughput datasets of system components (omics data), experimental methods of analysis and data integration, as well as the development and application of network approaches and computationally derived models. In pharmaceutical research, systems biology efforts are directed towards the identification of drug targets, the development of novel therapeutics and new indications for existing drugs. [Berg, 2014]

Omics tools, developed over the past several decades, can provide global information on the levels and dynamic changes in cellular and tissue components at specific time points in samples from cell-based assays, preclinical animal models or human studies. Omics data sets derived from transcriptomics, proteomics, and metabolomics are being used and integrated with each other, as well as with genomics information and other data types, to construct models of cell signaling, pathway and disease networks. These models can help to identify new targets as well as better understand and predict drug action *in vivo*.

The ultimate goal of systems biology in this context is an understanding of

physiology and disease across the multiple hierarchical levels of organization, from chemical and molecular interactions to pathways and pathway networks, at the cell-cell and tissue level, organs and organ systems and, ultimately, to the functioning of the whole organism. [Berg, 2014]

The data available for learning in discovery and development of pharmaceuticals tends to be compound-centric. Compound data are generally concerned with the identification and characterization of small molecules or biologics that selectively inhibit (or activate) specific molecular targets or pathway mechanisms. These kinds of studies are particularly useful to drug discovery research as the attempt to work out related drug mechanisms of action. They also support secondary drug development goals, such as clinical indication selection and patient stratification. [Berg, 2014]

1.7 Reproducibility

Several recent publications have highlighted the negative impact of irreproducible biomedical research, such as the group from Bayer that claimed they had to halt two-thirds of their research efforts in 2011 [Prinz et al., 2011] and the group from Amgen that reported only an 11% success rate in trying to replicate the effects of major cancer drug findings. [Begley and Lee, 2012] In the latter case we don't even know which drugs were tested because those findings are not public either. Information generation can happen far faster and is much more common than data analysis and knowledge creation in the biological sciences.

Some of the issues to overcome in this area include the need for more biologists trained in quantitative and statistical methods for analyzing large data sets and the open release of findings and experimental protocols.

Science conducted in an open fashion confers the following benefits

- Reproducibility of experiments allows other researchers to use the exact meth-

ods to calculate the relations between biological data.

- Faster development of disease models and therapeutic treatments due to the reuse of existing knowledge. Projects can be built upon existing results more easily or extend the research in directions unanticipated by the original team. First-pass results can be subject to new analysis and a second look at compounds with interesting side effects can lead to serendipitous discoveries.
- Increased quality as a result of having more researchers studying the same topic to provide a layer of assurance that errors will not propagate.
- Long-term availability of data and code. If these resources are not tied to businesses or patents, then they can be posted to multiple repositories to ensure that they are available in the future. [Prli and Lapp, 2012]

New research findings, supporting data and methods should, therefore, be made publicly available for independent verification and replication in order not to delay medical advances.

1.8 Aims

The aim of this project is to build QSAR models that quantify the risk of off-target effects for candidate molecules by building statistical machine learning models that make binary classification as inhibitors or non-inhibitors of Cytochrome P450 of compounds based on their structure.

We will compare a well accepted, commercial method of binary classification with a three open source implementations of QSAR model building according to the following plan:

- Build Binary QSAR models in the Molecular Operating Environment (MOE).
- Develop and implement comparable methods in open source software.
- Evaluate and compare results from all models.
- Perform this analysis as reproducible research.

Chapter 2

Materials and Methods

2.1 Review of PubChem Assay 1851

PubChem BioAssay 1851 contains data for inhibition of five major CYP isoforms (1A2, 2C9, 2C19, 2D6 and 3A4) by 17,143 chemical compounds. [Veith et al., 2009] The tested compounds were all drugs or drug-like compounds. A look at the chemical space occupied by these compounds reveals that the majority of them had molecular weight below 500 daltons and logP below 5. [Lapins et al., 2013]

PubChem Assay 1851 used low low-luminescent substrates, which are converted to more-luminescent metabolites as the signal for this assay. The progression of a reaction was measured by an increase in light intensity during CYP metabolism of the substrate. Inhibitors of a particular CYP isozyme reduced the rate of metabolism of the substrate and thus results in a decreased luminescent signal. [Zlokarnik et al., 2005]

The most recent technology developed for CYP inhibition is based on substrates that release luciferin as the metabolite. This is a coupled assay system in which addition of luciferase and ATP converts the freed luciferin to des-carboxyluciferin along with light emission. The format is similar to fluorescence methods, but particularly well-suited to high throughput applications because and all that is required are

addition-only manipulations and luminescence plate readers. [Zlokarnik et al., 2005] The assay obtained luminescence readings at a range of compound concentrations and then determined activity parameters using the Hill equation.

In the PubChem Assay 1851, compounds are classified as active or inactive inhibitors for each CYP with an activity cutoff set to $AC_{50} = 10\mu M$ (AC_{50} , activity concentration 50, refers to the concentration that is required to elicit half-maximal effect). However, in cases where the dose-response curve for a compound showed poor fit or the inhibition efficacy was below 60%, the assay results were regarded as inconclusive. [Lapins et al., 2013]

Compounds were characterized by their Activity Score and regarded as inhibitors if their activity score ranged between 40 and 100. An Activity Score is assigned based on AC_{50} value, which was combined with a confidence measure. Combining measures for completeness of a dose-response curve and efficacy of inhibition, resulted in the Activity Score, where a larger value indicates higher inhibitory activity and/or higher confidence in inhibitory assay result. Compounds with an activity score equal to zero are considered as non-inhibitors while compounds with activity scores above 0 and up to 40 are considered inconclusive. [Lapins et al., 2013]

The data from Pubchem Assay 1851 is available for download from the NIH through the PubChem website. The interface changes from time to time, but I downloaded two files which comprised the entire dataset for the experiment. The first file, a structures file, contained the structural information encoded in the simplified molecular-input line-entry system (SMILES) format for each tested compound with corresponding Structure ID (SID) and Compound ID (CID) as assigned by NCBI. Another file, also organized by SID and CID, contained all of the luminescent responses from the high-throughput screen and the fitted parameters which are then summarized by the Activity Score.

2.2 Data Set Preparation and Molecular Descriptor Generation

Both data files from Bioassay 1851 were downloaded as comma separated value (.csv) files and merged together based on the SID column using functions from Python’s pandas library, which is made for handling, manipulating and reshaping structured data.

The preliminary steps of data preprocessing typically require data cleaning because raw data often contain anomalies, errors, or inconsistencies such as missing data, incomplete data, and invalid character values which may cause trouble in data analysis if left untreated. It is more complicated when data are collated from many formats that require harmonization and redundancy elimination. [Nantasenamat, 2009] In this case, minimal preprocessing was required due to the polished nature of the source files.

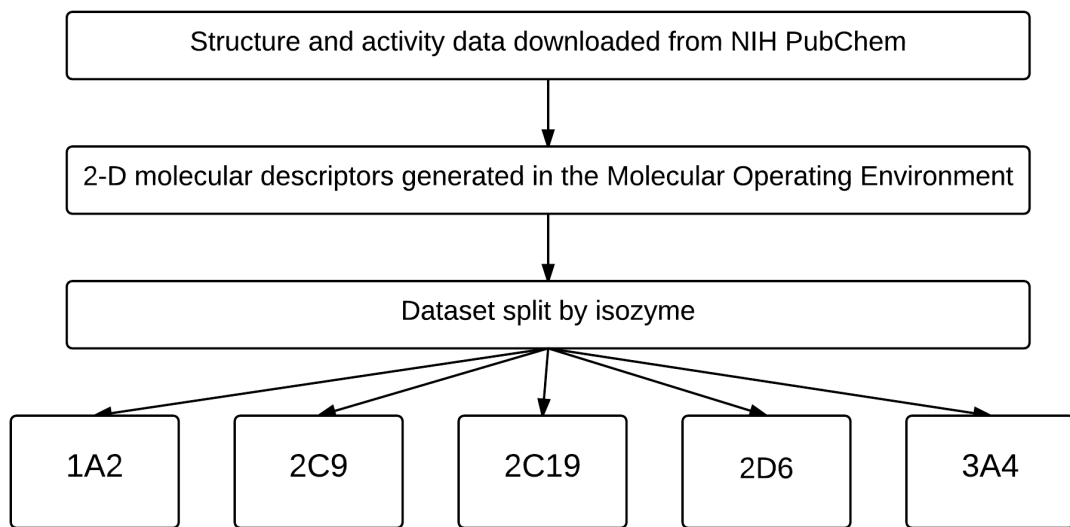


Figure 2.1: Data Download and Preparation

Molecular Descriptor Generation

First the merged file was loaded into a database in the Molecular Operating Environment. MOE functionality was used to obtain the washed configuration of compounds by removing the salts and finding an energy-minimized conformation. MOE was then used to calculate descriptors based on the molecular structures. The entire suite of MOE 2-D descriptors was selected for descriptor generation, resulting in 186 additional columns of nominal, ordinal and continuous values appended to the database. The resulting master file was saved in .csv format.

The dataset was then split by isozyme into 5 separate files that each contained only the SID, the Activity Score, and the 186 MOE 2-D descriptors using a script in the Python programming language.

For each isozyme, the number of active inhibitors was far outnumbered by the number of inactive inhibitors. The data file for each isozyme was subjected to a Python script that separated the inactives from actives. Then the order of the inactives was randomly shuffled and the column trimmed to the length of the activity column, thereby balancing the number of inactives and actives as required by some of the later statistical analyses.

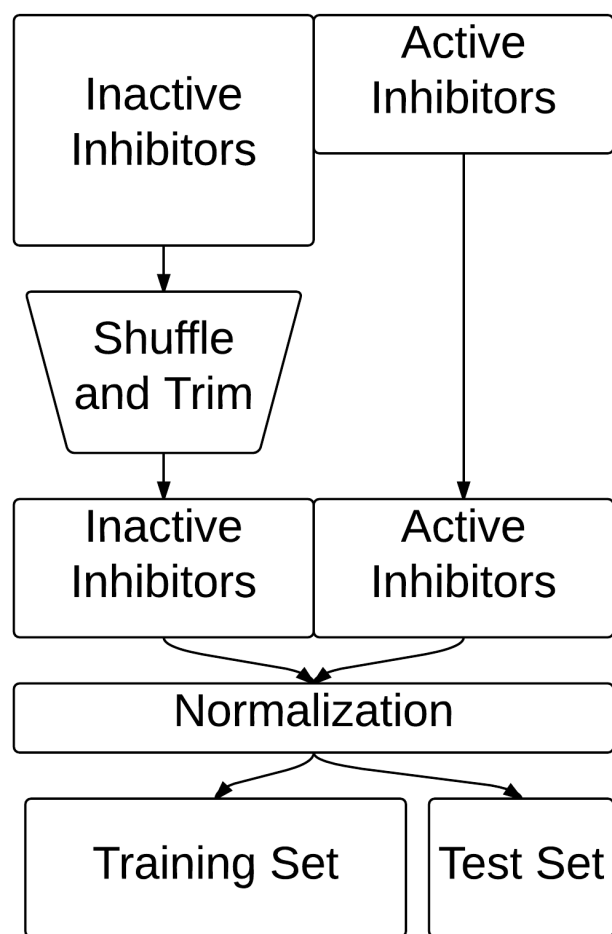


Figure 2.2: Isozyme Data Set Preparation

Next the script randomly shuffled the balanced datasets and split them into a training set and a test set. The training set held 80% of the original values and the test set held 20% of the original values. The split was not based on activity score. The ratio of active/inactive in each split was inspected to check if they were still acceptably balanced.

For each use of randomness in computation, a seed number was set for the random number generator to ensure reproducible results.

The balanced and split datasets are saved to Figshare.com for permanent, free and open access. (<http://dx.doi.org/10.6084/m9.figshare.1181846>) and

(<http://dx.doi.org/10.6084/m9.figshare.1066108>) All subsequent analyses use these same splits for comparability.

2.3 Feature selection

Typically data sets often contain redundant or noisy variables which make it more difficult for learning algorithms to discern meaningful patterns from the input. Such multicollinearity of the variables can either be used or treated if necessary to reduce computational resources required for model construction. [Nantasenamat, 2009]

There existed a great deal of variability in the range and distribution of each variable in the data set. This can pose problems for algorithms using distance measurements in the learning step. These situations are handled by applying statistical techniques such as min-max normalization or z-score standardiation.

In min-max normalization, the minimum and maximum value of each variable is adjusted to a uniform range between 0 and 1 according to the following equation: $X_{norm} = \frac{X - X_{min}}{X_{max} - X_{min}}$ In z-score standardization, essentially the variable of interest is subjected to statistical operation to achieve mean center and unit variance according to the following formula: $z = \frac{x - \mu}{\sigma}$ [Nantasenamat, 2009]

In situations where the data does not have a Gaussian (normal) distribution, simple mathematical functions can be applied to achieve normality or symmetry in the data distribution. A commonly used approach is to apply logarithmic transformation on the the variable of interest in order to achieve distribution approaching normality. This is typically performed on dependent variables such as the modeled biological/chemical properties of interest whereby IC50 may be transformed to logIC50 or -logIC50. Practically, such mathematical operations are applied to each individual value of a given variable of interest. [Nantasenamat, 2009]

2.4 Modeling - Mapping Descriptors to Activity Data

At this point a complete dataset is assembled for each isozyme that is balanced for the number of active and inactive inhibitors. That dataset is then further subdivided into a training set and a test set for later validation of results. The following section describes the different methods of feature selection, normalization and classification algorithms used in this study.

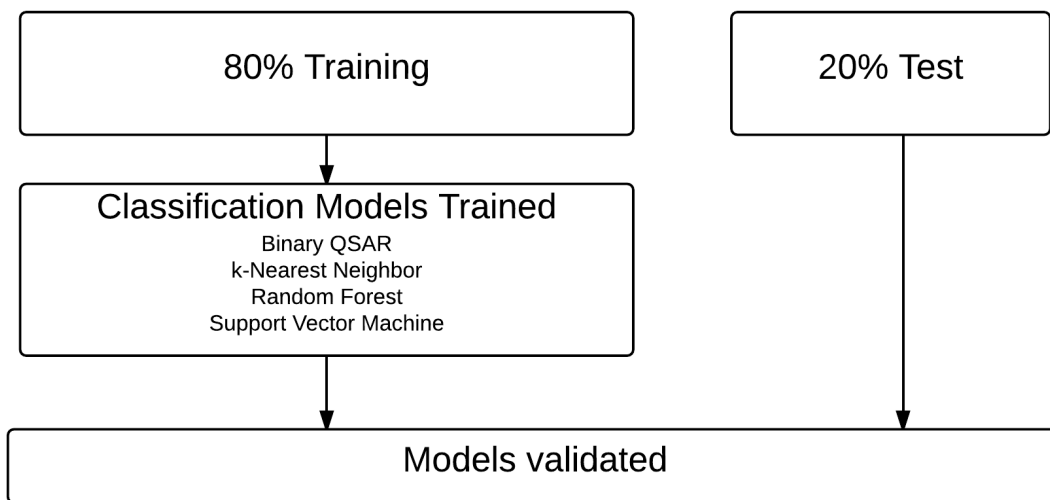


Figure 2.3: Supervised Classification Model Training and Validation

2.4.1 Binary QSAR in the Molecular Operating Environment

The Binary QSAR approach outlined by Labute [Labute, 1999] and included in the 2011 version of MOE, takes a Bayesian and probabilistic approach to classification of activity after reducing all the descriptors to principal components. The initial principal component analysis (PCA) of the descriptor measurements reduce the dimensionality of the dataset to a smaller set of dimensions where the new axes

represent the directions of greatest variance/spread of the data. The new axes are ranked according to the amount of variance each explains, with the first principal component accounting for the most variance, the second less variance and so on down the line. Models are subsequently built with the original data projected onto any or all of these principal components.

Initially, Partial Least Squares regression, a quantitative method available in the Molecular Operating Environment that also begins with data reduction by PCA, was tried in Dr. Zheng's lab but has shown poor predictive accuracy with this dataset as demonstrated by previous attempts in practice.

To carry out this analysis, the training data was loaded into MOE and then a menu driven interface was used to initiate the Binary QSAR method. A threshold Activity Score value of 39 was selected; meaning all activity values 40 and above were considered active and all values 39 and below were considered inactive. The smoothing parameter was left at the default value of 0.25. MOE automatically performs principal component analysis on high-dimensional datasets before Binary QSAR and there was no way to use the full suite of molecular descriptors untransformed.

For each isozyme a number of models were built, each using a different number of principal components during model building. Each principal component is orthogonal and uncorrelated to the rest, and so each one captures a portion of the total variance inherent in the dataset. The assumption is that inclusion of more principal components leads to more of the variance being accounted for in a classification decision. However, since each principal component is a linear combination of all the variables, the benefits of dimensionality reduction comes at the cost of interpretability. One result of this study will be to determine how many principal components are necessary to achieve good prediction results with the most dimensionality reduction .i.e the fewest number of principal components. For exploratory purposes, models with 2, 5, 10, 15, 20, 30, and 44 principal components were constructed.

MOE models were written to .fit files and the model report saved as a .txt file.

Modeling in Python

The Python programming language is a dynamically-typed, object-oriented, interpreted language. Its primary strength lies in the ease with which it allows a programmer to rapidly prototype a project, coupled with a powerful and mature set of standard libraries, like the scikit-learn package, that can facilitate large-scale production-level software engineering projects as well. Python has a very shallow learning curve and easily accessible, online learning materials. The following machine learning classifiers used in this study were implemented with Python's scikit-learn package.

2.4.2 κ -Nearest Neighbor (κ NN)

The κ NN algorithm predicts the class of a test set object based on the class membership of its κ most similar training set objects. [Lapins et al., 2013] κ -nearest neighbors algorithms find the κ -points that are closest to a point in question based on their attributes using a certain distance measure (e.g., Euclidean distance).

2.4.3 Random Forest (RF)

Random Forest is a classifier that consists of multiple decision trees. To borrow the language of graph theory, decision trees are made of nodes and branches. At each node, the dataset is split based on the value of some attribute that is selected so that the instances of different classes are predominately moved to different branches. Classification starts at the root node and is performed by passing the instances along the tree to leaf nodes. To introduce diversity between the trees of a random forest, a small subset of all features is randomly selected to take decisions at each node of

each tree. The classification decision is performed by considering results of all trees in a majority vote. [Lapins et al., 2013]

2.4.4 Support Vector Machines (SVM)

SVM is a machine learning technique for classification that tries to find the hyperplane that gives the greatest margin of separation between classes. SVMs apply a kernel function to each instance which projects the data into a higher-dimensional feature space before finding the hyperplane. [Lapins et al., 2013] The margin is defined as the minimum distance from data points to the hyperplane. The data points with the smallest margin are called support vectors and used to establish the final SVM classifier.

2.5 Model Evaluation and Validation

The building blocks of a successful QSAR model are the accuracy of the input data, selection of appropriate descriptors and statistical tools, and most importantly validation of the developed model. Validation is the process where reliability and relevance of a procedure are established for a predetermined purpose. For QSAR models validation should target robustness, prediction performances and applicability domain (AD) of the models. [Lapins et al., 2013]

Statistical evaluation in QSAR modeling is essential to validate the model as well as to evaluate its predictive performance. The predictive performance of a data set can be assessed by dividing it into a training set and a testing set. The training set is used for constructing a model and then the predictive performance of that model is evaluated on the testing set. Performance is typically assessed internally from the predictive performance of the training set, while external performance can be assessed from the predictive performance of the independent test set that has never been seen

by the training model.

For evaluation by the test set, two statistical metrics were used – the overall prediction Accuracy and the True Positive/Negative Rate. A commonly used approach for internal validation is known as N-fold cross-validation where a data set is partitioned into N number of folds. For example, in a 5-fold cross-validation 1 partition is left out as a testing set while the remaining 4 folds are used as the training set for model construction. And then validation is performed with the fold that was left out. In situations where the number of samples is limited, leave-one-out cross-validation is the preferred approach. In that case, the number of folds is equal to the number of samples present in the data set. [Nantasenamat, 2009]

In model building, 5-fold cross validation was used in conjunction with model building on the training set for all of the models built in Python. Once model building and selection was complete, the test set was used for external validation. Therefore we assessed the predictive ability of the models by performing cross-validation and external predictions.

Confusion Matrix

A confusion matrix (or error matrix or table of confusion) is a representational summary of the performance of a classifier. All data points used in model building are apportioned along one axis according to their actual class value, and along another axis according to the value predicted by that model. For binary classification this results in a square matrix with two 2 columns and 2 rows.

		Prediction Class	
		p	n
Actual Class	p'	True Positive	False Negative
	n'	False Positive	True Negative

In this study, results that are both actual inhibitors according to assay data and predicted inhibitors according to the model results are deemed True Positives (TP). Similarly, actual noninhibitors that are predicted to be so are called True Negatives (TN). In the case of actual inhibitors predicted to be noninhibitors, these are labelled False Negatives (FN) and are the result of Type I errors. Noninhibitors classified as inhibitors by the model are referred to as False Positives (FP) and are the result of Type II errors.

Accuracy

The accuracy metric provides general information about how many compounds are misclassified. Accuracy is simply the percentage of correctly classified instances over the total number of instances and is calculated as

$$ACC = \frac{(TP + TN)}{(TP + FP + TN + FN)}$$

where TP is the number of true positives, TN is the number of true negatives, FP is the number of false positives or over-predictions, and FN is the number of false negatives or missed predictions.

Accuracy is not an optimal measure of model performance if the data set is unbalanced (i.e. sizes of the classes are unequal) or if certain types of errors are to be

considered more serious than others (e.g. false negatives compared to false positives). [Lapins et al., 2013] For this reason, all datasets were rebalanced to contain roughly equal numbers of active and inactive inhibitors.

As a further interrogation of results, the true positive rate and the true negative rate are also calculated for all models. The true positive rate is the fraction of true positives out of all actual inhibitors and it is calculated as follows

$$TruePositiveRate = \frac{TP}{(TP + FN)}$$

The true positive rate is also referred to as Sensitivity and in MOE software it is called Accuracy on Active.

The true negative rate is the fraction of true negatives over the number of actual noninhibitors and is calculated as :

$$TrueNegativeRate = \frac{TN}{(TN + FP)}$$

The true negative rate is also known as the Specificity of a classification, and in MOE Binary QSAR it is called Accuracy on Inactives.

Binary QSAR after PLS of descriptors

The test sets were loaded and the washed structures were appended to the .csv files. All models were evaluated using the menu driven workflow in MOE and the classification probabilities were appended to the database file and saved as a .csv.

The resulting file was loaded into a Microsoft Excel spreadsheet and the predictions classified as actives or inactives. Predicted probabilities of ≥ 0.5 were evaluated as active inhibitor predictions and predicted probabilities of ≤ 0.5 were deemed inactive inhibitors. Confusion matrices were then tabulated within the spreadsheet for predicted vs actual actives and inactives. And from these confusion matrices, classi-

fication accuracy measures were calculated - total accuracy, TPR, and TNR. These results were saved and the results reported below.

κ -NN Utilizing the Full Set of 2-D Descriptors

The full models follow a similar procedure to the previous workflow, except they omit data reduction by principal component analysis and use the full descriptor set to construct and evaluate models.

For each isozyme separately, training data is loaded into a dataframe. The Activity Score is assigned the role of response variable and all 186 descriptors are identified as predictor variables. Predictor variables are scaled and normed to mean 0 and standard deviation of 1. `kNN.fit` method is called from the scikit-learn library to train a model, which then reports the confusion matrix and an accuracy score for the training model by five-fold cross-validation. TPR and TNR are calculated from the confusion matrix.

Random Forest Classification Utilizing the Full Set of 2-D Descriptors

For each isozyme separately, training data is loaded into a dataframe. The Activity Score is assigned the role of response variable and all 186 descriptors are identified as predictor variables. Predictor variables are scaled and normed to mean 0 and standard deviation of 1. `RF.fit` method is called from the scikit-learn library to train a model, which then reports the confusion matrix and an accuracy score for the training model by five-fold cross-validation. TPR and TNR are calculated from the confusion matrix.

Support Vector Machine classification utilizing the full set of 2-D descriptors

For each isozyme separately, training data is loaded into a dataframe. The Activity Score is assigned the role of response variable and all 186 descriptors are identified

as predictor variables. Predictor variables are scaled and normed to mean 0 and standard deviation of 1. SVM.fit method is called from the scikit-learn library to train a model, which then reports a the confusion matrix and an accuracy score for the training model by five-fold cross-validation. TPR and TNR are calculated from the confusion matrix.

A script was written for each isozyme that performed these three fit methods in series. Code and results are documented in IPython notebooks that are currently hosted on github.com (<https://github.com/5x5x5x5/CYP/>) and freely accessible and downloadable. Presented in this way, they are more easily verifiable and extendable. Experimenting with other classification algorithms in scikit-learn simply requires adding new method.fit calls to the model building loop, because of the simple and consistent design of the scikit-learn library.

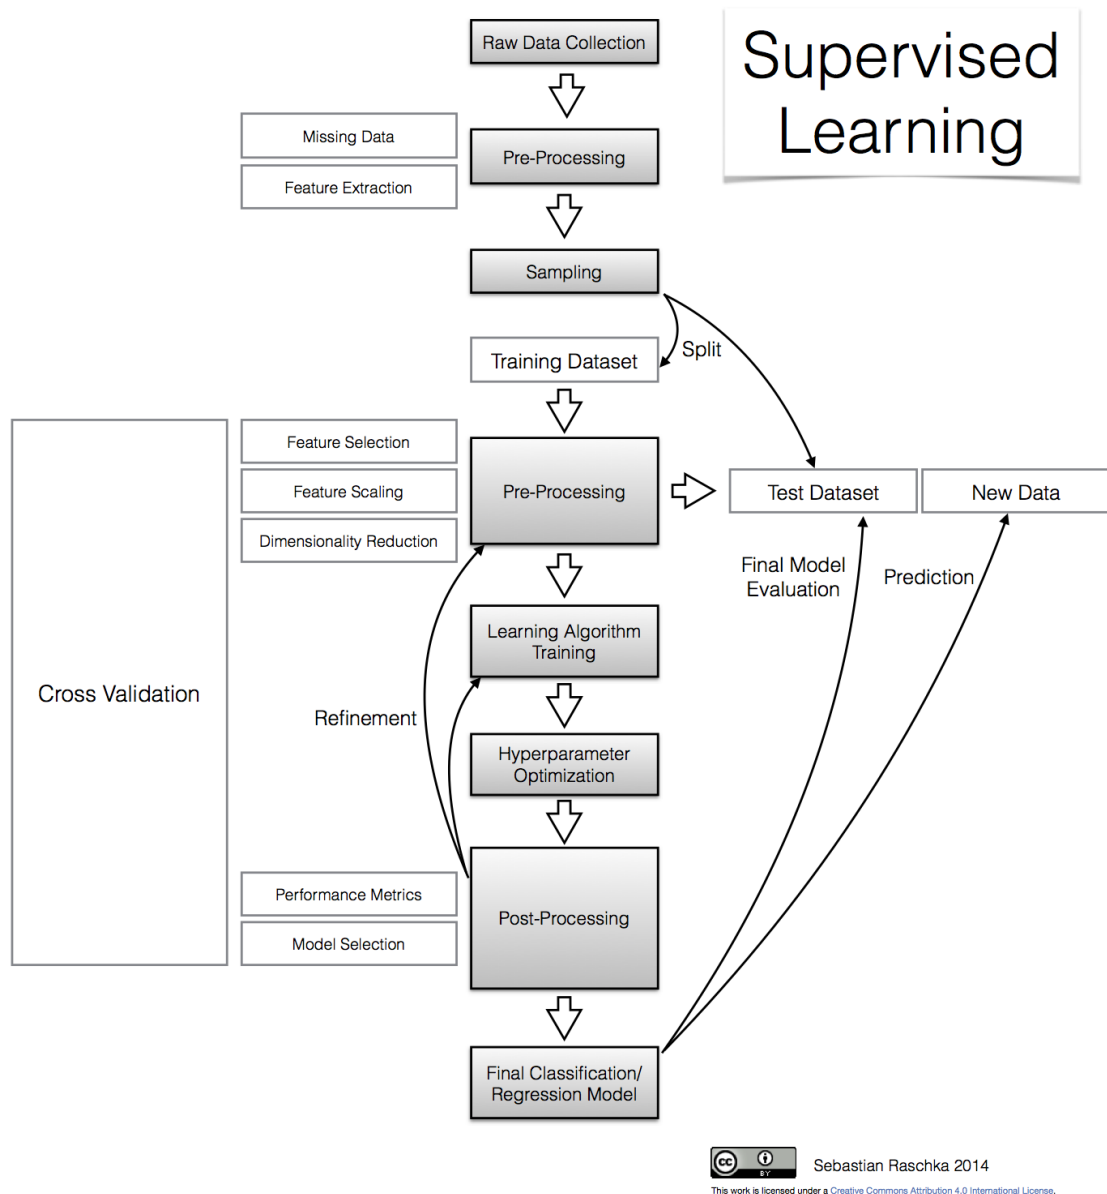


Figure 2.4: Supervised Learning Workflow for Classification

Chapter 3

Results

3.1 1A2 MOE Models

Table 3.1: 1A2 MOE Model Results

1A2 Training Set Accuracy				1A2 Test Set Accuracy			
PC	Total	Active	Inactive	PC	Total	Active	Inactive
2	0.638	0.752	0.522	2	0.632	0.758	0.517
5	0.704	0.746	0.662	5	0.705	0.773	0.643
10	0.734	0.755	0.712	10	0.746	0.749	0.743
15	0.737	0.770	0.704	15	0.752	0.775	0.731
20	0.741	0.781	0.701	20	0.748	0.779	0.721
30	0.735	0.782	0.686	30	0.739	0.783	0.698
44	0.725	0.777	0.672	44	0.720	0.773	0.672

Cytochrome P450 isozyme 1A2 models were trained on 4820 actives and 4780 inactives in the training set, and validated against 529 actives and 580 inactives in the test set. All models are considered good models when the total accuracy is above 0.6. The test set results represent compounds unseen by the model. Validating these models against the test set does not produce wildly divergent accuracy scores, indicating that overfitting is unlikely.




Figure 3.1: 1A2 MOE Model Accuracy

A visual scan of Figure 3.1 shows that model accuracy is higher for predicting active compounds than inactive compounds, and total accuracy is somewhere in between. 20 principal components (PC20) appears to be an optimal number for capturing relevant variance. PC20 is approximately the point at which total accuracy is maximized and dropoff of accuracy on inactives at either end of the range is avoided. At 20PCs, the total accuracy on the test set is 0.748. Accuracy on Actives is 0.779 and accuracy on inactives is 0.721.

3.2 2C9 MOE Models

Table 3.2: 2C9 MOE Model Results

2C9 Training Set Accuracy			
PC	Total	Active	Inactive
2	0.620	0.695	0.547
5	0.683	0.722	0.645
10	0.699	0.744	0.655
15	0.702	0.752	0.653
20	0.710	0.756	0.665
30	0.711	0.754	0.668
44	0.712	0.770	0.655

2C9 Test Set Accuracy			
PC	Total	Active	Inactive
2	0.609	0.702	0.510
5	0.675	0.713	0.634
10	0.682	0.730	0.632
15	0.686	0.739	0.630
20	0.685	0.731	0.636
30	0.690	0.730	0.646
44	0.688	0.750	0.621

Cytochrome P450 isozyme 2C9 models were trained on 3266 actives and 3327 inactives in the training set, and validated against 855 actives and 794 inactives in the test set. All models are considered good models when the total accuracy is above 0.6. The test set results represent compounds unseen by the model. Validating these models against the test set does not produce wildly divergent accuracy scores, indicating that overfitting is unlikely.




Figure 3.2: 2C19 MOE Model Accuracy

A visual scan of Figure 3.2 shows that model accuracy is higher for predicting active compounds than inactive compounds, and total accuracy is somewhere in between. 20 principal components appears to be an optimal number for capturing relevant variance. PC20 is approximately the point at which total accuracy is maximized and dropoff of accuracy on inactives at either end of the range is avoided. At 20PCs, the total accuracy on the test set is 0.685. Accuracy on Actives is 0.731 and accuracy on inactives is 0.636.

3.3 2C19 MOE models

Table 3.3: 2C19 MOE Model Results

2C19 Training Set Accuracy				2C19 Test Set Accuracy			
PC	Total	Active	Inactive	PC	Total	Active	Inactive
2	0.585	0.774	0.396	2	0.593	0.775	0.408
5	0.685	0.726	0.645	5	0.683	0.729	0.635
10	0.699	0.736	0.661	10	0.691	0.731	0.650
15	0.700	0.747	0.653	15	0.687	0.736	0.637
20	0.705	0.749	0.661	20	0.698	0.758	0.638
30	0.717	0.761	0.672	30	0.699	0.764	0.633
44	0.708	0.777	0.640	44	0.694	0.774	0.613

Cytochrome P450 isozyme 2C19 models were trained on 4721 actives and 4741 inactives in the training set, and validated against 1193 actives and 1173 inactives in the test set. All models are considered good models when the total accuracy is above 0.6. The test set results represent compounds unseen by the model. Validating these models against the test set does not produce wildly divergent accuracy scores, indicating that overfitting is unlikely.




Figure 3.3: 2C19 MOE Model Accuracy

A visual scan of Figure 3.3 shows that model accuracy is higher for predicting

active compounds than inactive compounds, and total accuracy is somewhere in between. 20 principal components appears to be an optimal number for capturing relevant variance. PC20 is approximately the point at which total accuracy is maximized and dropoff of accuracy on inactives at either end of the range is avoided. At 20PCs, the total accuracy on the test set is 0.698. Accuracy on Actives is 0.758 and accuracy on inactives is 0.638.

3.4 2D6 MOE Models

Table 3.4: 2D6 MOE Model Results

2D6 Training Set Accuracy				2D6 Test Set Accuracy			
PC	Total	Active	Inactive	PC	Total	Active	Inactive
2	0.589	0.753	0.426	2	0.590	0.738	0.443
5	0.670	0.706	0.634	5	0.667	0.685	0.650
10	0.685	0.721	0.648	10	0.662	0.688	0.636
15	0.686	0.709	0.662	15	0.665	0.692	0.639
20	0.705	0.721	0.689	20	0.683	0.699	0.668
30	0.703	0.732	0.673	30	0.686	0.696	0.677
44	0.690	0.725	0.656	44	0.669	0.707	0.632

Cytochrome P450 isozyme 2D6 models were trained on 2219 of actives and 2214 of inactives in the training set, and validated against 552 actives and 557 inactives in the test set. All models are considered good models when the total accuracy is above 0.6. The test set results represent compounds unseen by the model. Validating these models against the test set does not produce wildly divergent accuracy scores, indicating that overfitting is unlikely.




Figure 3.4: 2D6 MOE Model Accuracy

A visual scan of Figure 3.4 shows that model accuracy is higher for predicting active compounds than inactive compounds, and total accuracy is somewhere in between. 20 principal components appears to be an optimal number for capturing relevant variance. PC20 is approximately the point at which total accuracy is maximized and dropoff of accuracy on inactives at either end of the range is avoided. At 20 PCs, the total accuracy on the test set is 0.683. Accuracy on Actives is 0.699 and accuracy on inactives is 0.668.

3.5 3A4 MOE Models

Table 3.5: 3A4 MOE Model Results

3A4 Training Set Accuracy				3A4 Test Set Accuracy			
PC	Total	Active	Inactive	PC	Total	Active	Inactive
2	0.644	0.699	0.589	2	0.627	0.683	0.573
5	0.675	0.727	0.623	5	0.656	0.719	0.595
10	0.698	0.741	0.655	10	0.677	0.719	0.637
15	0.701	0.746	0.656	15	0.680	0.733	0.628
20	0.705	0.749	0.660	20	0.686	0.736	0.638
30	0.709	0.761	0.657	30	0.686	0.742	0.631
44	0.698	0.773	0.623	44	0.676	0.747	0.607

Cytochrome P450 isozyme 3A4 models were trained on 4234 of actives and 4193 of inactives in the training set, and validated against 1033 actives and 1074 inactives in the test set. All models are considered good models when the total accuracy is above 0.6. The test set results represent compounds unseen by the model. Validating these models against the test set does not produce wildly divergent accuracy scores, indicating that overfitting is unlikely.




Figure 3.5: 3A4 MOE Model Accuracy

A visual scan of Figure 3.5 shows that model accuracy is higher for predicting active compounds than inactive compounds, and total accuracy is somewhere in between. 20 principal components appears to be an optimal number for capturing relevant variance. PC20 is approximately the point at which total accuracy is maximized and dropoff of accuracy on inactives at either end of the range is avoided. At 20PCs, the total accuracy on the test set is 0.686. Accuracy on Actives is 0.736 and accuracy on inactives is 0.638.

3.6 1A2 Classification Method Comparison

Table 3.6: Comparison of Classification Methods for 1A2

1A2 Classification Method Comparison		
	Training Set	Test Set
MOE 20PC	0.741	0.748
kNN	0.761	0.754
Random Forest	0.770	0.769
SVD	0.800	0.804

The comparison of model accuracy results on the CYP 1A2 isozyme dataset show that Binary QSAR implemented in the Molecular Operating Environment gave an overall accuracy of 0.748 for the test set, while the methods implemented by Python’s scikit-learn gave accuracies of 0.754 for k-nearest neighbor, 0.769 for random forests and 0.804 for support vector machines on the test set.

3.7 2C9 Classification Method Comparison

Table 3.7: Comparison of Classification Methods for 2C9

2C9 Classification Method Comparison		
	Training Set	Test Set
MOE 20PC	0.710	0.685
kNN	0.721	0.692
Random Forest	0.717	0.687
SVD	0.749	0.720

The comparison of model accuracy results on the CYP 2C9 isozyme dataset show that Binary QSAR implemented in the Molecular Operating Environment gave an overall accuracy of 0.748 for the test set, while the methods implemented by Python’s scikit-learn gave accuracies of 0.692 for k-nearest neighbor, 0.687 for random forests and 0.720 for support vector machines on the test set.

3.8 2C19 Classification Method Comparison

Table 3.8: Comparison of Classification Methods for 2C19

2C19 Classification Method Comparison		
	Training Set	Test Set
MOE 20PC	0.705	0.698
kNN	0.730	0.720
Random Forest	0.736	0.721
SVD	0.767	0.756

The comparison of model accuracy results on the CYP 2C19 isozyme dataset show that Binary QSAR implemented in the Molecular Operating Environment gave an overall accuracy of 0.698 for the test set, while the methods implemented by Python’s scikit-learn gave accuracies of 0.720 for k-nearest neighbor, 0.721 for random forests and 0.756 for support vector machines on the test set.

3.9 2D6 Classification Method Comparison

Table 3.9: Comparison of Classification Methods for 2D6

2D6 Classification Method Comparison		
	Training Set	Test Set
MOE 20PC	0.705	0.683
kNN	0.717	0.678
Random Forest	0.729	0.707
SVD	0.755	0.725

The comparison of model accuracy results on the CYP 2D6 isozyme dataset show that Binary QSAR implemented in the Molecular Operating Environment gave an overall accuracy of 0.683 for the test set, while the methods implemented by Python’s scikit-learn gave accuracies of 0.678 for k-nearest neighbor, 0.707 for random forests and 0.725 for support vector machines on the test set.

3.10 3A4 Classification Method Comparison

Table 3.10: Comparison of Classification Methods for 3A4

3A4 Classification Method Comparison		
	Training Set	Test Set
MOE 20PC	0.705	0.686
kNN	0.716	0.674
Random Forest	0.719	0.683
SVD	0.755	0.714

The comparison of model accuracy results on the CYP 3A4 isozyme dataset show that Binary QSAR implemented in the Molecular Operating Environment gave an overall accuracy of 0.686 for the test set, while the methods implemented by Python's scikit-learn gave accuracies of 0.674 for k-nearest neighbor, 0.683 for random forests and 0.714 for support vector machines on the test set.

3.11 Overall Method Comparison

Inhibitor Classification Method Comparison - Training Set




Figure 3.6: Inhibitor Classification Method Comparison on the Training Set

Looking at the results for method comparison on the test set, it is clear that all methods performed similarly for each isozyme and well overall. It appears that the support vector machine classifier slightly outperforms all others in each case.

Inhibitor Classification Method Comparison on the Test Set

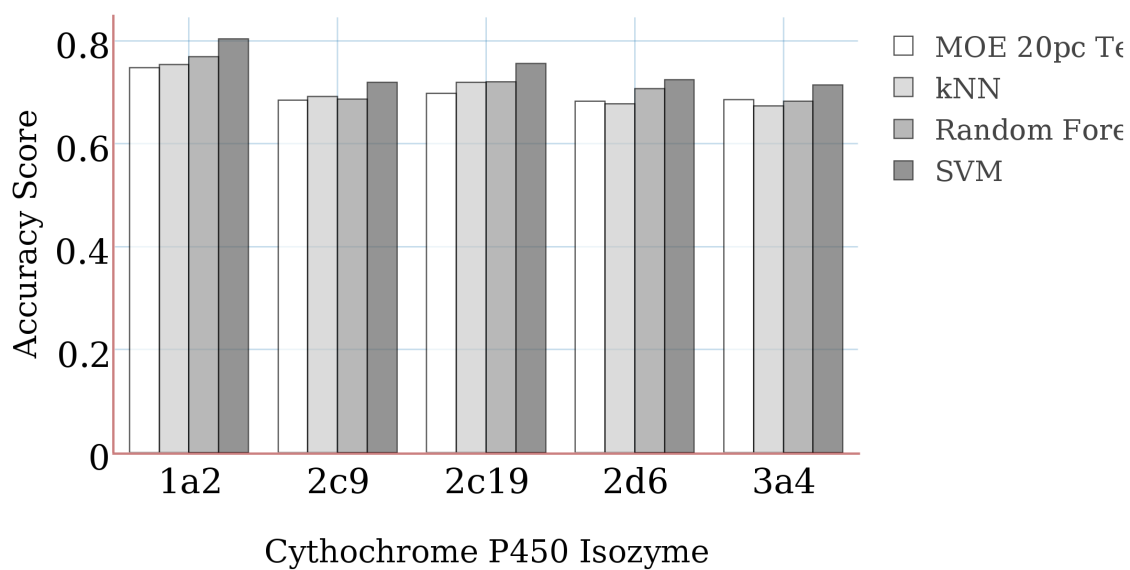


Figure 3.7: Inhibitor Classification Method Comparison on the Test Set

Looking at the results for method comparison on the test set, it is clear that all methods performed similarly for each isozyme and well overall. It appears that the support vector machine classifier slightly outperforms all others in each case.

Chapter 4

Discussion

4.1 Other Attempts at Modeling Assay 1851

Pubchem Bioassay AID1851 has been the basis for several attempts to advance *in silico* screening development. Several research groups have developed large-scale single target models for single isoforms. [Vasanthanathan et al., 2009] [Sridhar et al., 2012] [Sun, 2012] [Novotarskyi et al., 2010] Lapins et al. pursued all five isozymes but used a proteochemometric method that also takes into account each isozyme’s protein sequence. [Lapins et al., 2013] Cheng et al. created single target QSAR models for all five CYP isoforms and these are the most comparable to my efforts. [Cheng et al., 2011] These models all showed good predictive performances according to our standards of a 0.6 accuracy threshold.

In Cheng’s study inhibitor predicting models were developed for five major CYP isoforms, namely 1A2, 2C9, 2C19, 2D6, and 3A4. They used a combined classifier algorithm on a large data set containing more than 24,700 unique compounds, that came from PubChem and included the 17,143 used in this study. His group used an ensemble of different independent machine learning classifiers including support vector machine, C 4.5 decision tree, κ -nearest neighbor, and naive Bayes. The ensemble was

brought together by a back-propagation artificial neural network (BP-ANN). Those models were also validated by 5-fold cross-validation and tested on a test set composed of about 9000 diverse unique compounds previously unseen by the classifier. The range of the validation measures from the test set results was from 0.764 to 0.886. Cheng et al. claim these classification models are applicable for virtual screening of the five major CYP isoforms inhibitors or can be used as simple filters of potential chemicals in drug discovery. [Cheng et al., 2011] These results are similar to the findings presented above.

4.2 Sources of Error

This assay required recombinant sources of CYP enzymes, because the probes are not sufficiently P450 isoform-selective to be used with those derived from human liver microsomes. Another potential source for false positives in these assays can be compounds that interfere with light generation directly, such as compounds that interfere with luciferase enzymatic activity. [Zlokarnik et al., 2005]

The results of the large-scale screening of Pubchem AID1851 against five CYP isoforms identified that the majority of compounds in a typical chemical library cross-inhibited several isoforms, while only a small fraction of the compounds did not inhibit any of the isoforms. [Veith et al., 2009]

The phenomenon of uniformly lower accuracy on the inactives versus the actives might be explained by the composition of the inactive class. Both inactive inhibitors as well as compounds of uncertain activity were both included under the label 'inactive' for binary classification purposes. The chosen representation of class inclusion may have less potential for successful discrimination than. To test assumptions about class labels, the number of classes could be increased or the Activity Score threshold could be altered in future models.

4.3 Machine Learning for Pharmaceutical Sciences

Developing successful machine learning applications still requires a substantial amount of black art that is hard to find in textbooks. [Domingos, 2012]

Unlike in this study, raw data is often not in a form that is amenable to learning, but you can construct features from it that are. Easily the most important factor in a good model is the features used. If there are many independent features that correlate well with the class, learning is easy. However if the class is a very complex function of the features, the time required to train a good model may take too long or be impossible. [Domingos, 2012] Most of the time and effort in a machine learning project typically goes into feature engineering, and it is one area where domain expertise and experience with chemical library design can mean the difference between success and failure.

Also machine learning is not a one-shot process of building a data set and running a learner, but rather an iterative process of running the learner, analyzing results, modifying the data and/or learner, and repeating. Learning is often the quickest part of this, but only because of all the effort that has gone into sharing and codifying that knowledge. Feature engineering is more difficult because it is domain-specific, while learners can by and large be general-purpose. [Domingos, 2012]

There is ultimately no replacement for domain expertise in feature engineering when the other option is to run a learner with a very large number of features to find out which ones are useful in combination. The latter approach may be too time-consuming, or cause overfitting. [Eklund et al., 2014]

As a rule, it pays to try the simplest learners first (e.g. naive Bayes before logistic regression, κ -nearest neighbor before support vector machines). More sophisticated learners may be seductive, but they are usually harder to use, because there are more tuning parameters and their internals are more opaque. [Domingos, 2012]

One way to divide learning algorithms is as follows: those whose representation has a fixed size, like linear classifiers, and those whose representation can grow with the data, like decision trees. Fixed-size learners can only take advantage of so much data. Variable-sized learners can in principle learn any function given sufficient data, but in practice may not, because of limitations of the algorithm or computational cost. Also, because of the curse of dimensionality, no existing amount of data may be enough. [Domingos, 2012] Each type of learner comes with its own assumptions and, as of yet, none are demonstrably best in all situations, although some are clearly better than most. [Hand, 2006, Fernández-Delgado et al., 2014]

Perhaps because each type of learner is looking at the data from different angles, combining the results of different learners into an ensemble of models is a technique that has demonstrated its effectiveness at improving overall accuracy. Creating model ensembles is now standard in machine learning. The Statistics community uses techniques, such as bagging, boosting and stacking to resample training data or reweight classifier inputs on previously misclassified data in order to reduce variance while minimizing bias. But the simplest forms of ensemble of modeling can be thought of as a 'majority vote', where final class assignment of an instance rests at the class arrived at by most of the included models.

Ensembles of models frequently outperform any single model, so a logical next step of this study would be to combine results from all models developed.

Chapter 5

Conclusion

QSAR models and their value in informing regulatory decision making will likely increase with the standardization of analytical approaches, more complete and reliable data collection methods and a better understanding of toxicity mechanisms in the role of disease as well as individual susceptibility to adverse clinical events. [Kruhlak et al., 2012]

I have demonstrated several methods for building models for binary classification of compounds for inhibition of cytochrome P450 isozymes. They all do pretty well and are significantly better than randomly guessing. There are many more machine learning methods to try, of particular interest are neural networks and more Bayesian methods. The data and code for model comparison for this study is available online, so there are no artificial barriers to benchmarking further results.

Bibliography

- [Begley and Lee, 2012] Begley, C. G. and Lee, E. M. (2012). Drug development: Raise standards for preclinical cancer research. *Nature*, 10:531–533.
- [Berg, 2014] Berg, E. L. (2014). Systems biology in drug discovery and development. *Drug Discovery Today*, 19(2):113–25.
- [Boulesteix and Schmid, 2014] Boulesteix, A. L. and Schmid, M. (2014). Machine Learning Versus Statistical Modeling. *Biometrical Journal*, 4(56):588–593.
- [Breiman, 2001] Breiman, L. (2001). Statistical Modeling : The Two Cultures. *Statistical Science*, 16(3):199–231.
- [Cheng et al., 2011] Cheng, F., Yu, Y., Shen, J., Yang, L., Li, W., Liu, G., Lee, P. W., and Tang, Y. (2011). Classification of cytochrome P450 inhibitors and noninhibitors using combined classifiers. *Journal of Chemical Information and Modeling*, 51(5):996–1011.
- [Danielson, 2002] Danielson, P. B. (2002). The cytochrome P450 superfamily: biochemistry, evolution and drug metabolism in humans. *Current Drug Metabolism*, 3(6):561–97.
- [Domingos, 2012] Domingos, P. (2012). A few useful things to know about machine learning. *Communications of the ACM*, 55(10):78.
- [Eklund et al., 2014] Eklund, M., Norinder, U., Boyer, S., and Carlsson, L. (2014). Choosing Feature Selection and Learning Algorithms in QSAR. *Journal of Chemical Information and Modeling*, 54(3):837–43.
- [Fernández-Delgado et al., 2014] Fernández-Delgado, M., Cernadas, E., Barro, S., and Amorim, D. (2014). Do we need hundreds of classifiers to solve real world classification problems? *Journal of Machine Learning Research*, 15:3133–3181.
- [Hand, 2006] Hand, D. J. (2006). Classifier technology and the illusion of progress. *Statistical Science*, 21(1):1–14.
- [Hansch and Fujita, 1964] Hansch, C. and Fujita, T. (1964). A Method for the Correlation of Biological Activity and Chemical Structure. *Journal of American Chemical Society*, 86:1616–1626.

- [James et al., 2013] James, G., Witten, D., Hastie, T., and Tibshirani, R. (2013). *An Introduction to Statistical Learning*. Springer.
- [Kruhlak et al., 2012] Kruhlak, N. L., Benz, R. D., Zhou, H., and Colatsky, T. J. (2012). (Q)SAR modeling and safety assessment in regulatory review. *Clinical Pharmacology and Therapeutics*, 91:529–34.
- [Labute, 1999] Labute, P. (1999). Binary qsar: A new method for determination of quantitative structure activity relationships. *Pacific Symposium on Biocomputing*, 4:444–455.
- [Lapins et al., 2013] Lapins, M., Worachartcheewan, A., Spjuth, O., Georgiev, V., Prachayasittikul, V., Nantasenamat, C., and Wikberg, J. E. S. (2013). A unified proteochemometric model for prediction of inhibition of cytochrome p450 isoforms. *PLoS One*, 8(6):e66566.
- [Nantasenamat, 2009] Nantasenamat, C. (2009). A Practical Overview of Quantitative Structure-Activity Relationship. *EXCLI J*, 8:74–88.
- [Novotarskyi et al., 2010] Novotarskyi, S., Sushko, I., Körner, R., Pandey, A., and Tetko, I. (2010). Classification of CYP450 1A2 Inhibitors Using PubChem Data. *Journal of Cheminformatics*, 2(Suppl 1):P40.
- [Prinz et al., 2011] Prinz, F., Schlange, T., and Asadullah, K. (2011). Believe it or not: How much can we rely on published data on potential drug targets? *Nat Rev Drug Discovery*, 10:712.
- [Prli and Lapp, 2012] Prli, A. and Lapp, H. (2012). Standards for the plos computational biology software section. *PLoS Comput Biol*, 8(11):e1002799.
- [Scannell et al., 2012] Scannell, J. W., Blanckley, A., Boldon, H., and Warrington, B. (2012). Diagnosing the decline in pharmaceutical R&D efficiency. *Nature Reviews. Drug discovery*, 11(3):191–200.
- [Singh et al., 2011] Singh, D., Kashyap, A., Pandey, R., and Saini, K. (2011). Novel advances in cytochrome P450 research. *Drug Discovery Today*, 16(17-18):793–9.
- [Sridhar et al., 2012] Sridhar, J., Liu, J., Foroozesh, M., and Stevens, C. L. K. (2012). Insights on cytochrome p450 enzymes and inhibitors obtained through QSAR studies. *Molecules (Basel, Switzerland)*, 17(8):9283–305.
- [Sun, 2012] Sun, H. (2012). Predictive Models for Cytochrome P450 Isozymes Based on Quantitative High Throughput Screening Data. *J. Chem Inf Model*, 51(10):2474–2481.
- [Vasanthanathan et al., 2009] Vasanthanathan, P., Taboureau, O., Oostenbrink, C., Vermeulen, N. P., Olsen, L., and Jørgensen, F. (2009). Classification of Cytochrome P450 1A2 Inhibitors and Noninhibitors by Machine Learning Techniques. *Drug Metab Dispos.*, 37(3):658–664.

- [Veith et al., 2009] Veith, H., Southall, N., Huang, R., and James, T. (2009). Comprehensive characterization of cytochrome P450 isozyme selectivity across chemical libraries. *Nature*, 27(11):1050–1055.
- [Visser et al., 2014] Visser, S. G., de Alwis, D. P., Kerbusch, T., Stone, J. a., and Allerheiligen, S. R. B. (2014). Implementation of Quantitative and Systems Pharmacology in Large Pharma. *CPT: Pharmacometrics & Systems Pharmacology*, 3(October):e142.
- [Zlokarnik et al., 2005] Zlokarnik, G., Grootenhuis, P., and Watson, J. (2005). High throughput P450 inhibition screens in early drug discovery . *Drug Discovery Today*, 10(21):1443–1450.